

Digital Image Processing (CSE/ECE 478)

Lecture # 10: Edges and Hough Transform

Avinash Sharma

Center for Visual Information Technology (CVIT),
IIIT Hyderabad

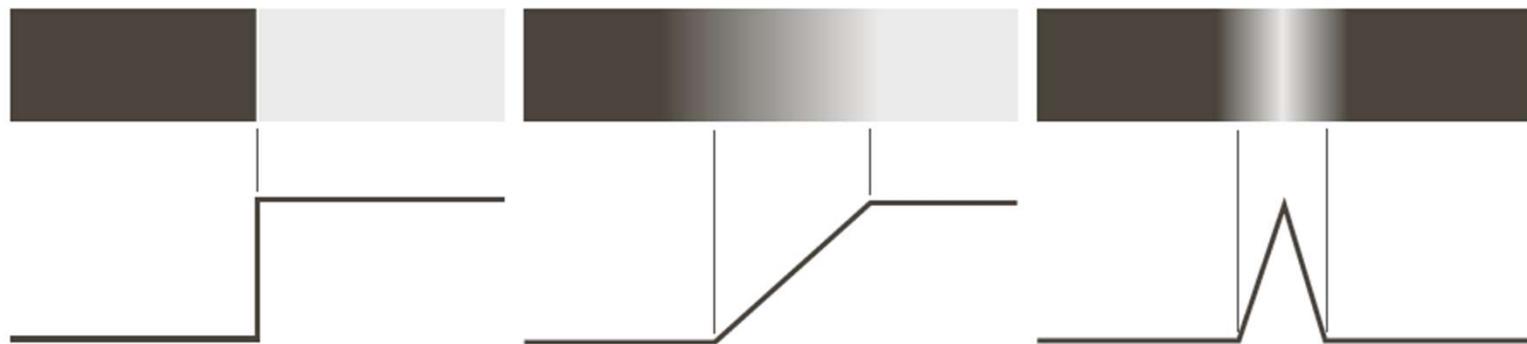


Today's Lecture

- Edge models, detection
- Advanced edge detection methods
- HOG (Histogram of Oriented Gradients) features
- Hough Transform



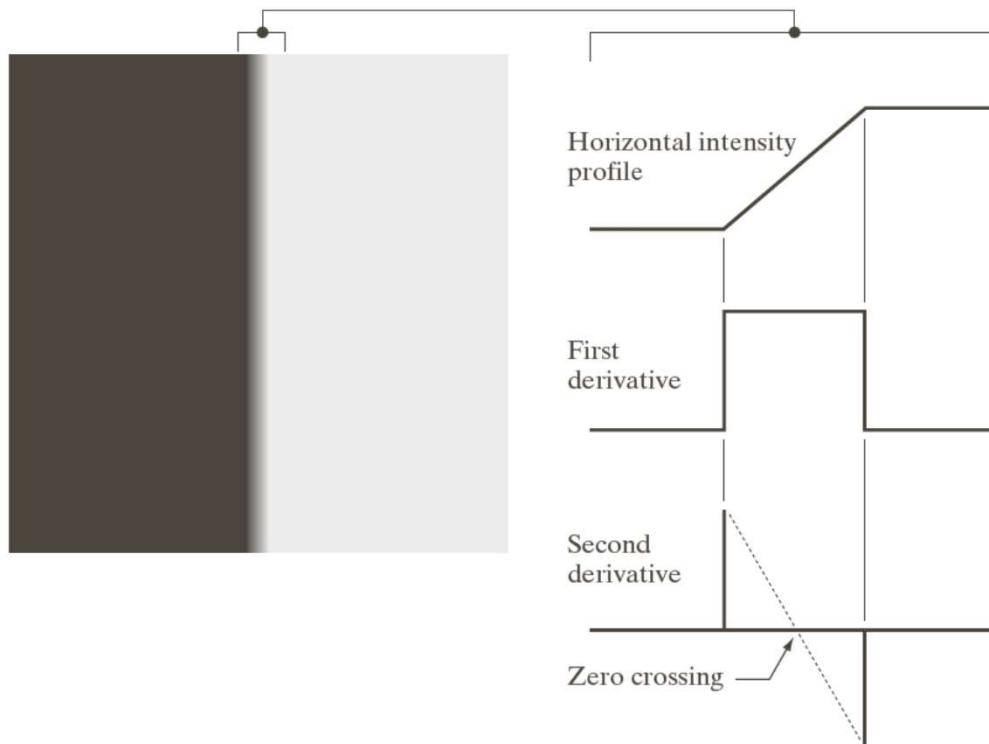
Edge Models



a b c

FIGURE 10.8
From left to right,
models (ideal
representations) of
a step, a ramp, and
a roof edge, and
their corresponding
intensity profiles.

Edge Models



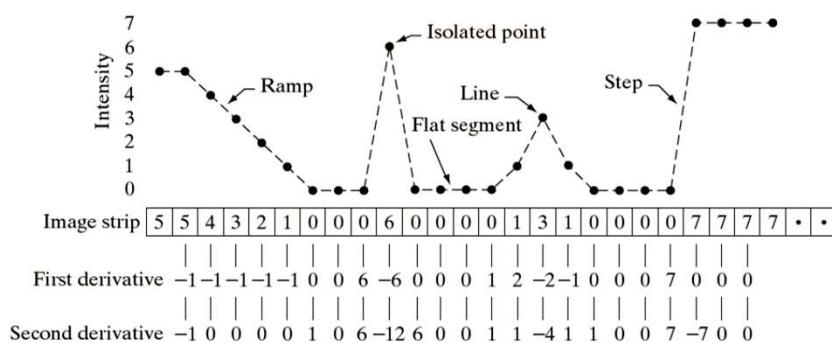
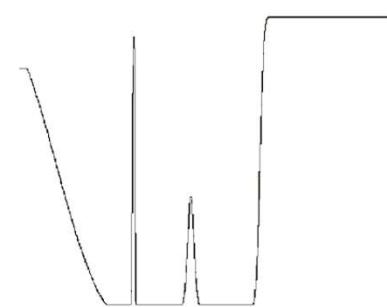
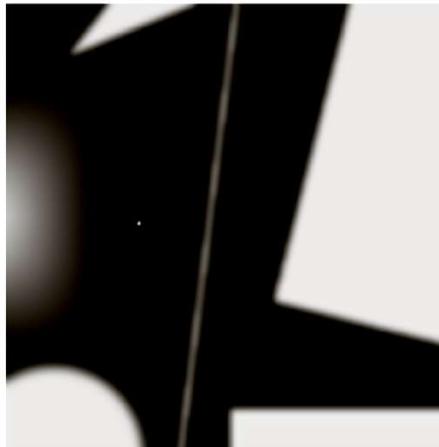
Discrete difference as Edge Detection

- Local changes in intensity → derivatives

	First order derivative	Second order derivative
Constant intensity	zero	zero
Onset of step or ramp	non zero	non zero
Along ramp	non zero	zero

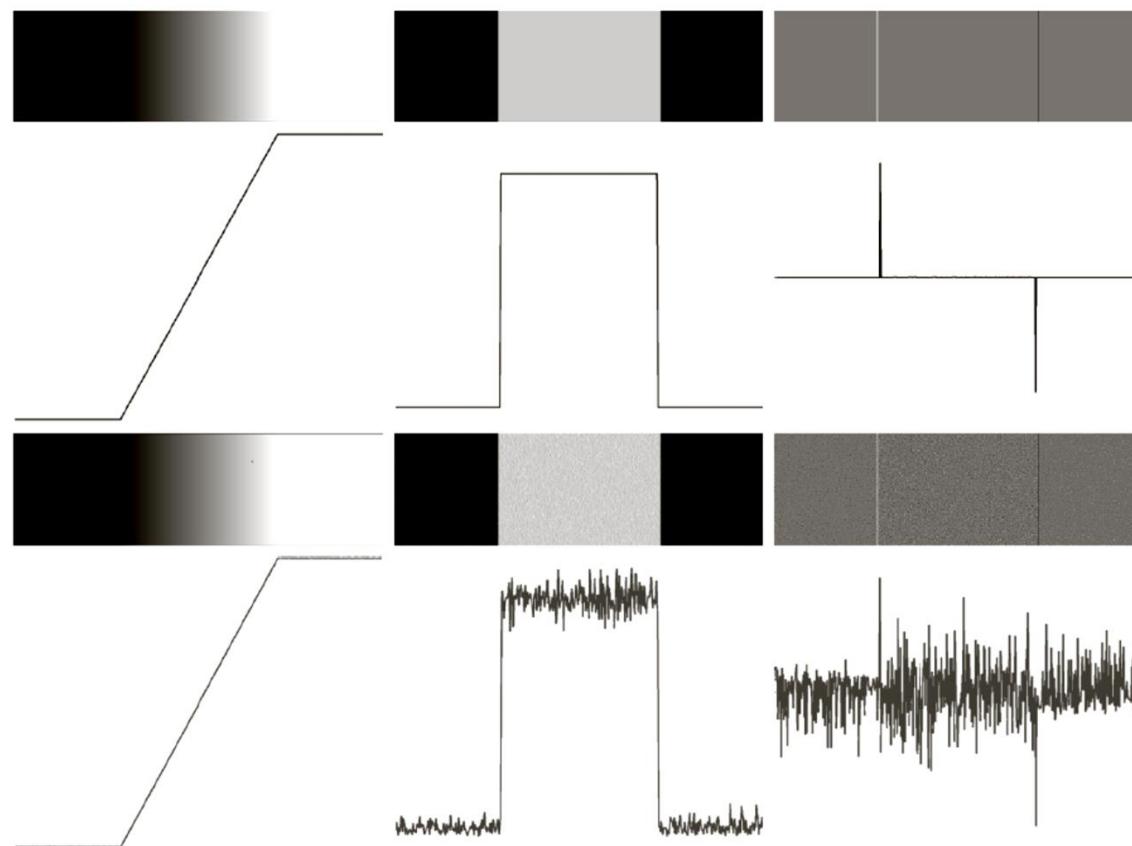


Edge Detection

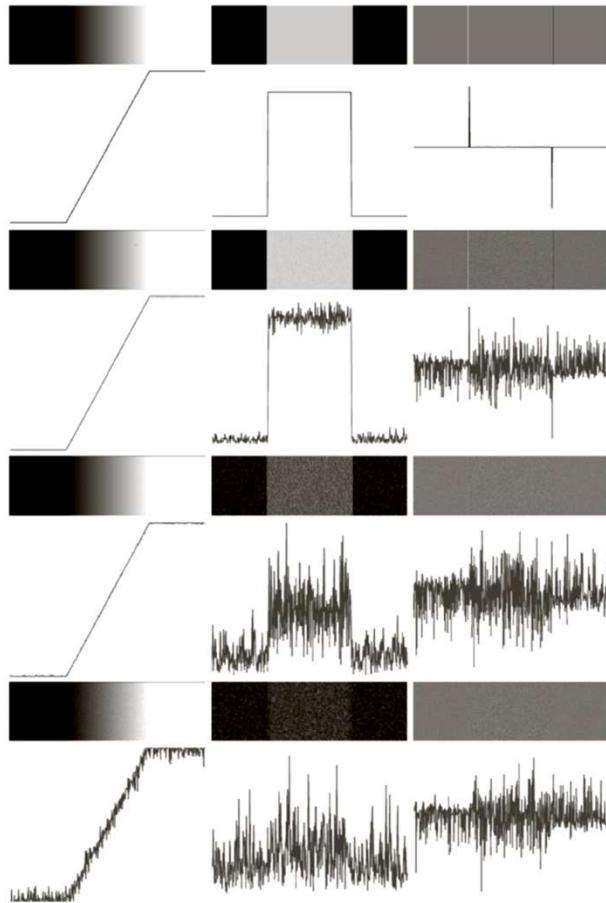


- 1st order derivative generally produces thick edges
- 2nd order derivative has a stronger response to fine details
- 2nd order derivative produce a double edge response
- The sign of 2nd order derivative can be used to determine whether the edge is a transition from light to dark or dark to light

Edge Models (effect of noise)



Edge Models (effect of noise)



Gaussian noise of zero mean
and standard deviation 0.0, 0.1,
1.0 and 10.0

Lets explore some more robust edge detectors!

The image gradient and its properties

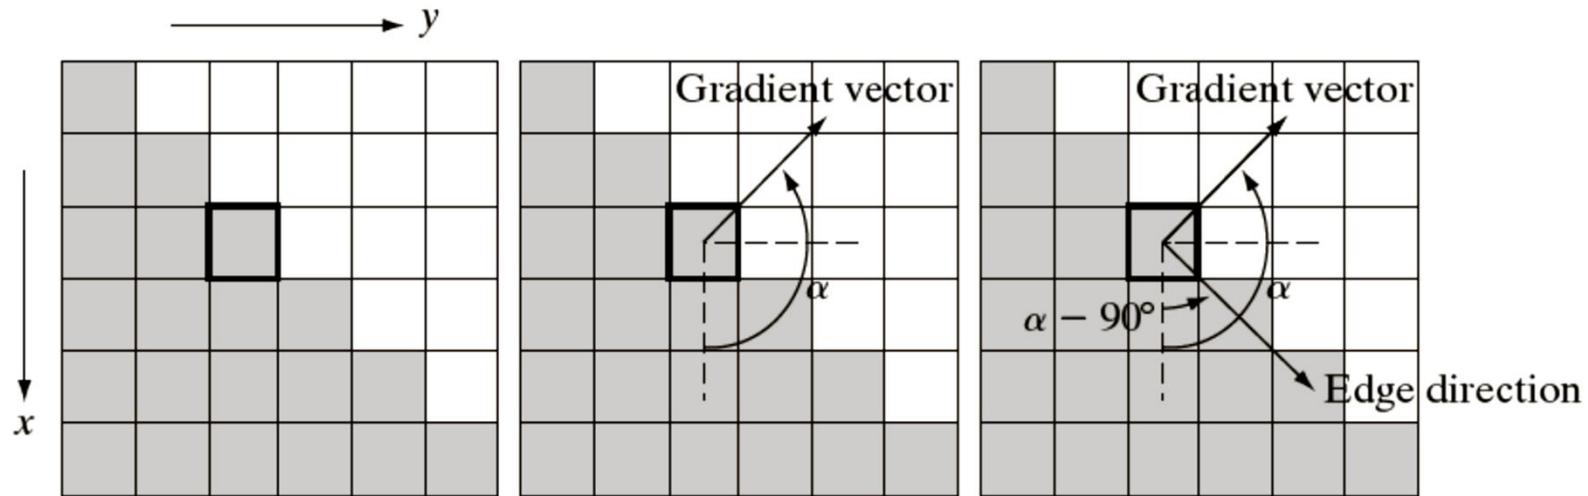
$$\nabla f = \text{grad}(f) = \begin{bmatrix} g_x \\ g_y \end{bmatrix} = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$$

$$M(x, y) = \text{mag}(\nabla f) = \sqrt{{g_x}^2 + {g_y}^2}$$

$$\alpha(x, y) = \tan^{-1} \left(\frac{g_x}{g_y} \right)$$



The image gradient and its properties



$$\nabla f = \text{grad}(f) = \begin{bmatrix} g_x \\ g_y \end{bmatrix} = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$$

$$M(x, y) = \text{mag}(\nabla f) = \sqrt{g_x^2 + g_y^2}$$

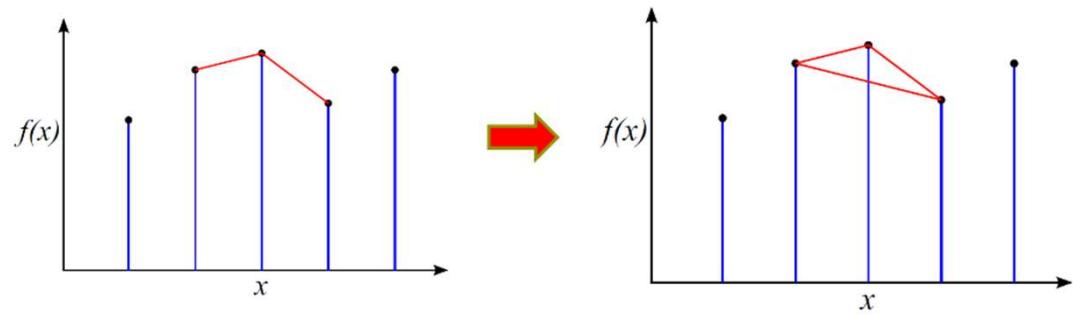
$$\alpha(x, y) = \tan^{-1} \left(\frac{g_x}{g_y} \right)$$



Gradient operators

$$\begin{bmatrix} -1 \\ 1 \end{bmatrix}$$

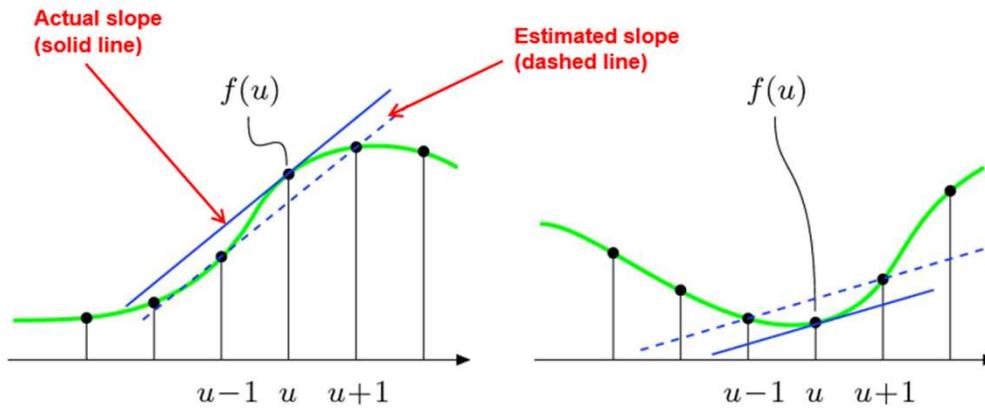
$$\begin{bmatrix} -1 & 1 \end{bmatrix}$$



- Left and right slope may not be same
- Solution?
Take average of left and right slope



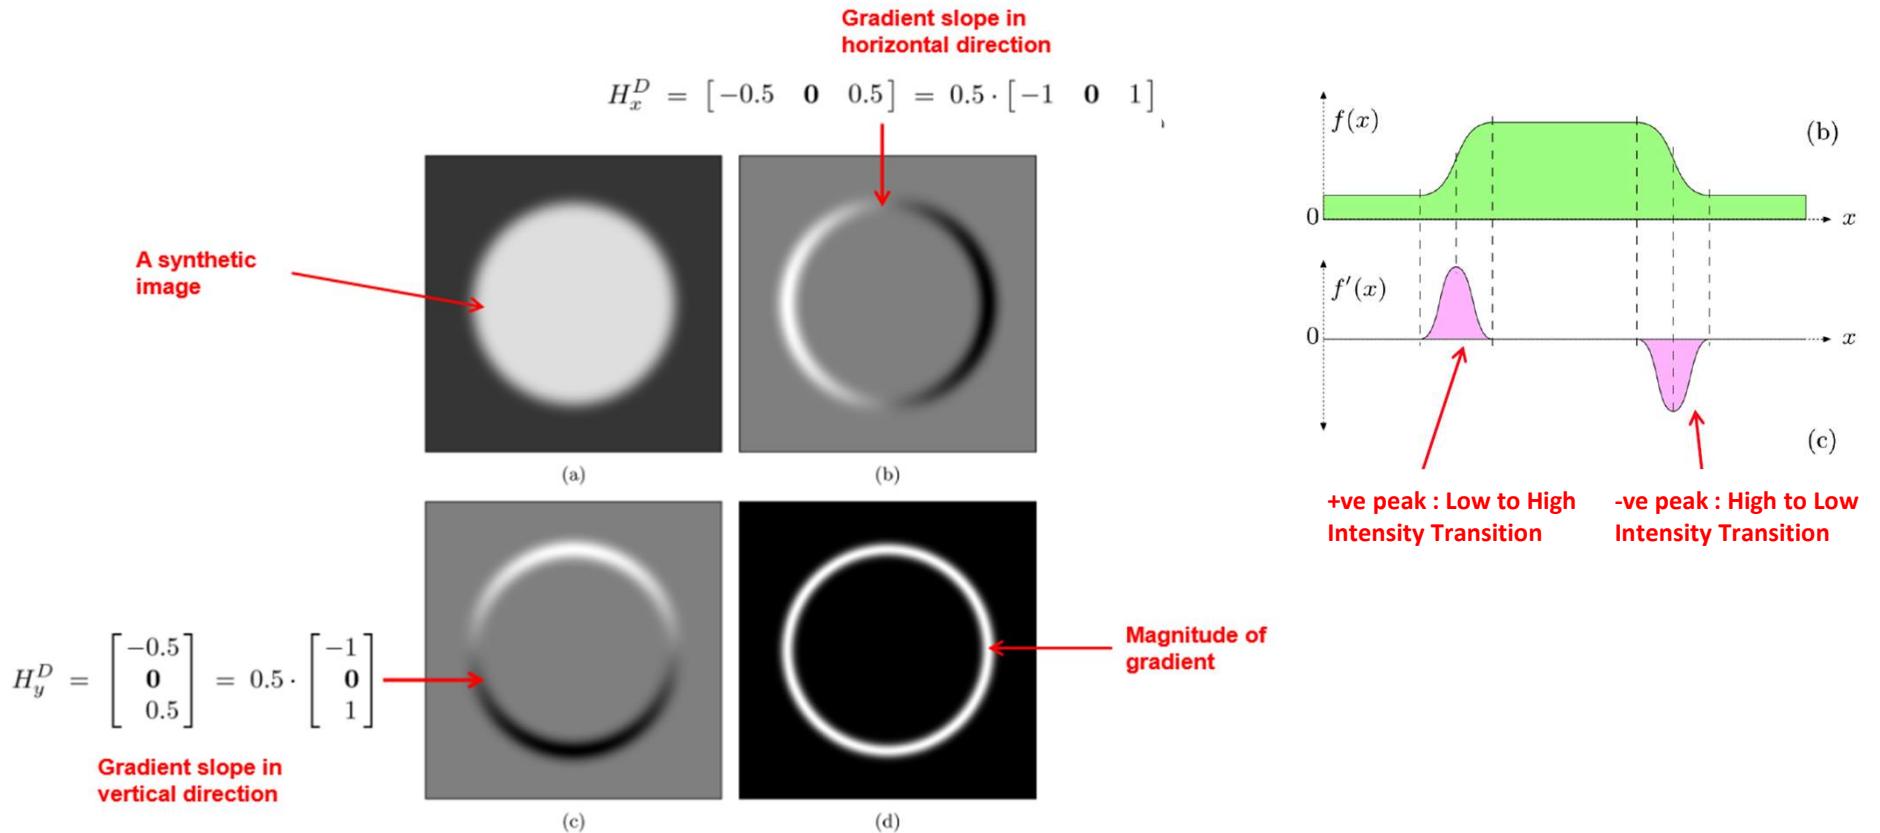
Gradient operators



$$\frac{df}{du}(u) \approx \frac{f(u+1) - f(u-1)}{2} = 0.5 \cdot (f(u+1) - f(u-1))$$



Gradient operators



Source: Emmanuel Agu

Common Gradient operators

-1
1

-1	1
----	---

-1	-1	-1
0	0	0
1	1	1

Prewitt

-1	0
0	1

0	-1
1	0

Roberts

-1	-2	-1
0	0	0
1	2	1

Sobel

- Symmetric filters take into account the nature of the data on opposite sides of the center point

Gradient operators (Prewitt)

- Finite differences sensitive to noise
- Derivative more robust if derivative computations are averaged in an neighbourhood

$$H_x^P = \frac{1}{3} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} * [0.5 \ 0 \ -0.5] = \frac{1}{6} \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$

Average in y direction

Derivative in x direction

-1	-1	-1
0	0	0
1	1	1

-1	0	1
-1	0	1
-1	0	1

Prewitt

Separability!

Gradient operators (Sobel)

$$H_x^S = \frac{1}{4} \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * [0.5 \ 0 \ -0.5] = \frac{1}{8} \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$H_y^S = \frac{1}{4} [1 \ 2 \ 1] * \begin{bmatrix} 0.5 \\ 0 \\ -0.5 \end{bmatrix} = \frac{1}{8} \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Average in x direction

Derivative in y direction

-1	-2	-1	-1	0	1
0	0	0	-2	0	2
1	2	1	-1	0	1

Sobel

Gradient operators

0	1	1
-1	0	1
-1	-1	0

-1	-1	0
-1	0	1
0	1	1

Prewitt

0	1	2
-1	0	1
-2	-1	0

-2	-1	0
-1	0	1
0	1	2

Sobel



Gradient operators



a b
c d

FIGURE 10.16
(a) Original image of size 834×1114 pixels, with intensity values scaled to the range $[0, 1]$.
(b) $|g_x|$, the component of the gradient in the x -direction, obtained using the Sobel mask in Fig. 10.14(f) to filter the image.
(c) $|g_y|$, obtained using the mask in Fig. 10.14(g).
(d) The gradient image, $|g_x| + |g_y|$.

Gradient operators

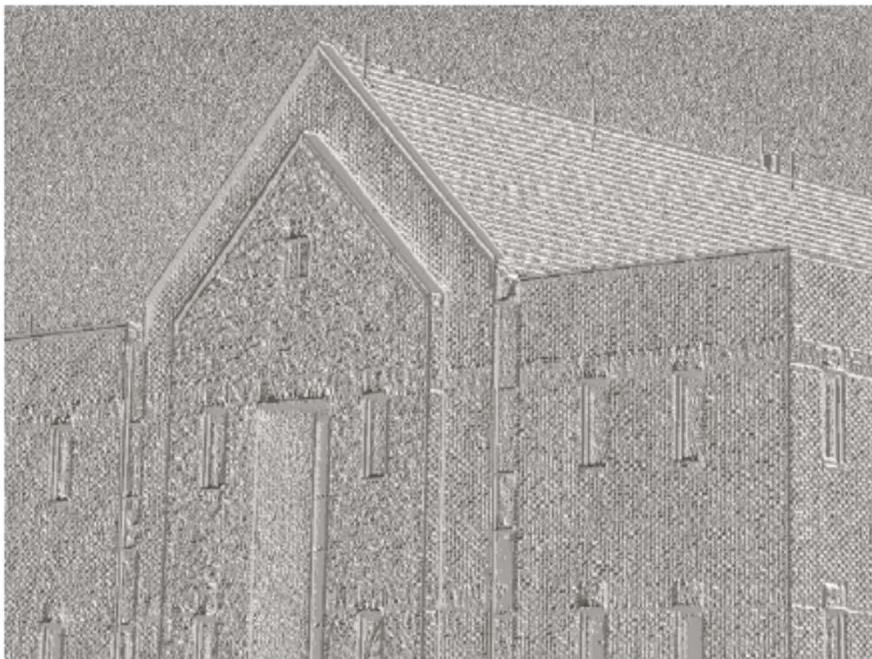
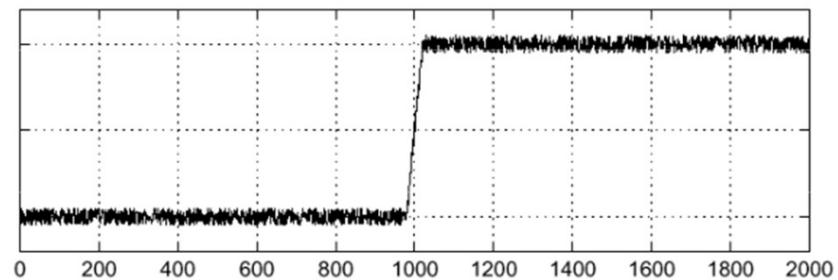


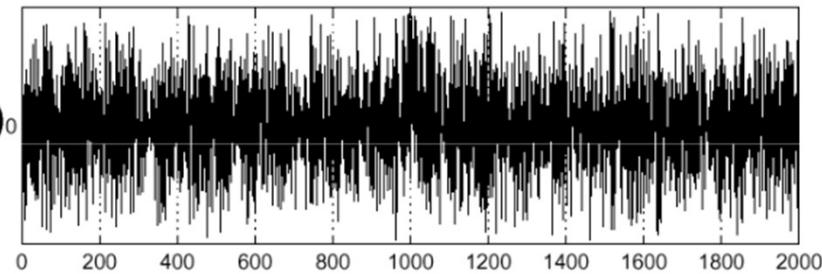
FIGURE 10.17
Gradient angle
image computed
using
Eq. (10.2-11).
Areas of constant
intensity in this
image indicate
that the direction
of the gradient
vector is the same
at all the pixel
locations in those
regions.

Where is the edge?

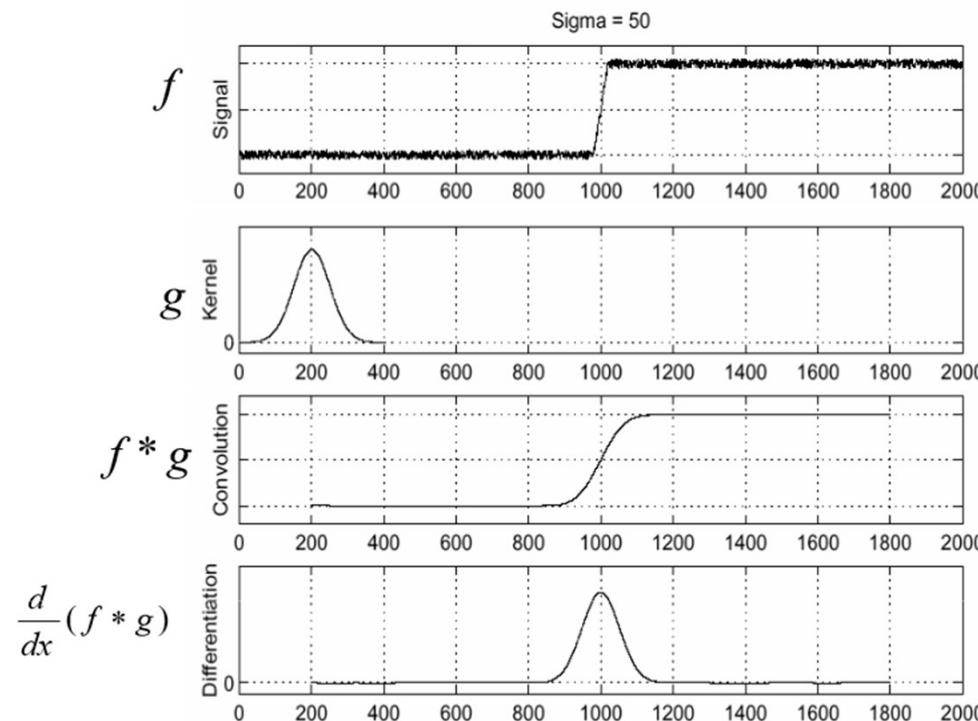
$f(x)$



$\frac{d}{dx}f(x)_0$

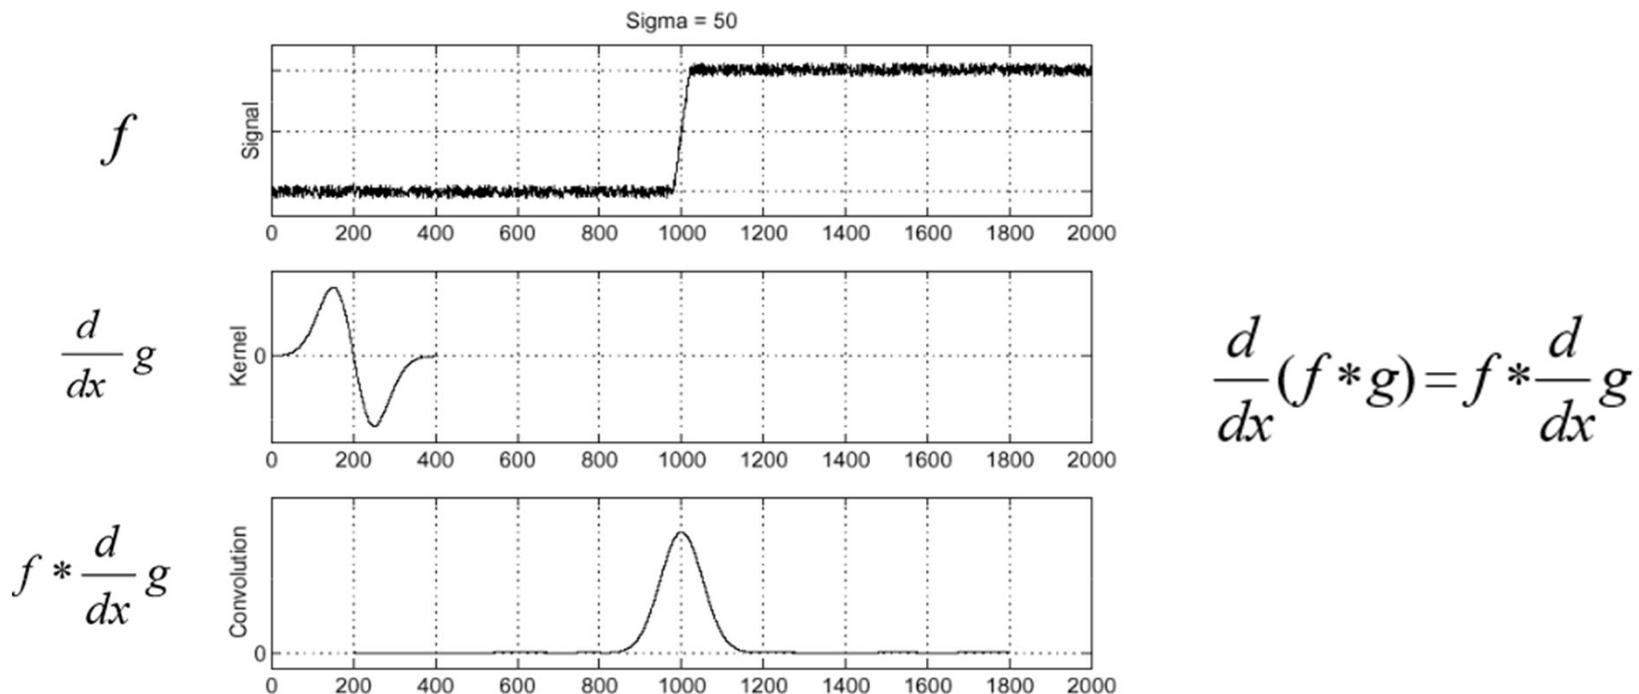


Solution: smooth first



Can do it faster!

- Differentiation in convolution and convolution is associative



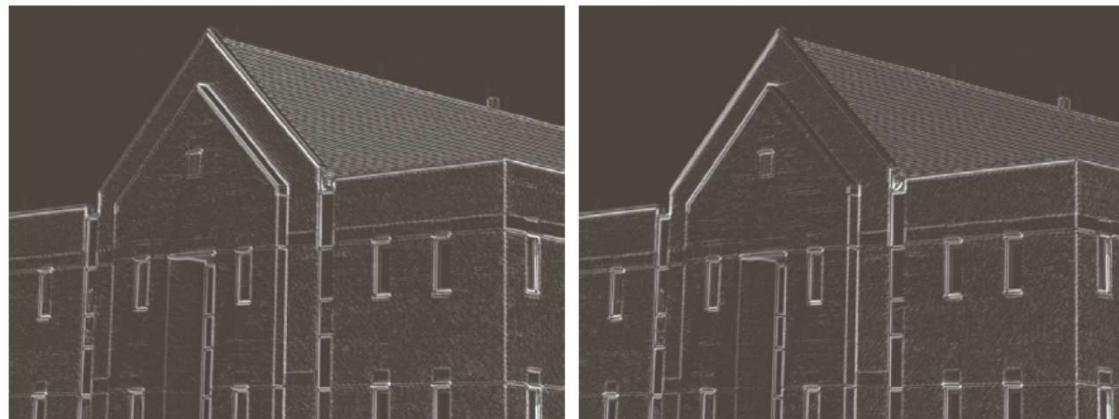
Gradient operators (with smoothing)



a b
c d

FIGURE 10.18
Same sequence as in Fig. 10.16, but with the original image smoothed using a 5×5 averaging filter prior to edge detection.

Gradient operators



a b

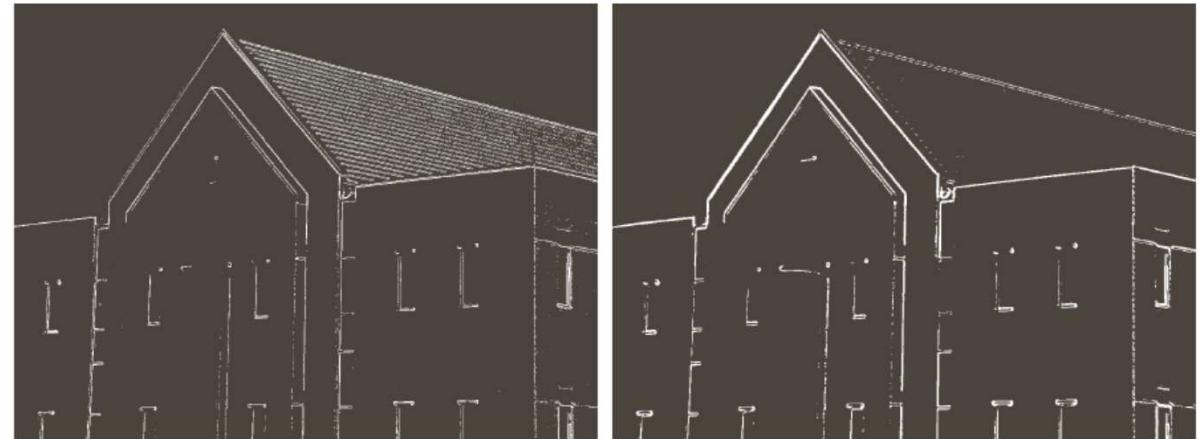
FIGURE 10.19
Diagonal edge detection.
(a) Result of using the mask in Fig. 10.15(c).
(b) Result of using the mask in Fig. 10.15(d). The input image in both cases was Fig. 10.18(a).

0	1	2
-1	0	1
-2	-1	0

-2	-1	0
-1	0	1
0	1	2

Sobel

Gradient + Thresholding



a b

FIGURE 10.20 (a) Thresholded version of the image in Fig. 10.16(d), with the threshold selected as 33% of the highest value in the image; this threshold was just high enough to eliminate most of the brick edges in the gradient image. (b) Thresholded version of the image in Fig. 10.18(d), obtained using a threshold equal to 33% of the highest value in that image.

Gradient + Thresholding



Original Image



Gradient magnitude, using Sobel operator

```
[fx, fy] = gradient(I);  
fmag = (fx.^2 + fy.^2) .^ .5;
```

Image courtesy: William Hoff

Gradient + Thresholding

- We can threshold edge magnitudes to produce binary image



Low threshold



High threshold



Non maximal suppression

Image courtesy: William Hoff

How much to smooth?

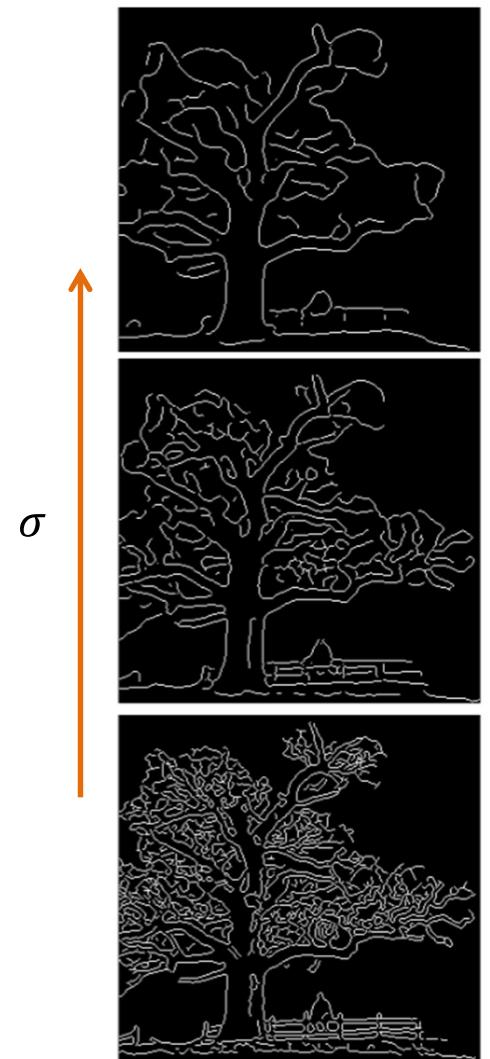
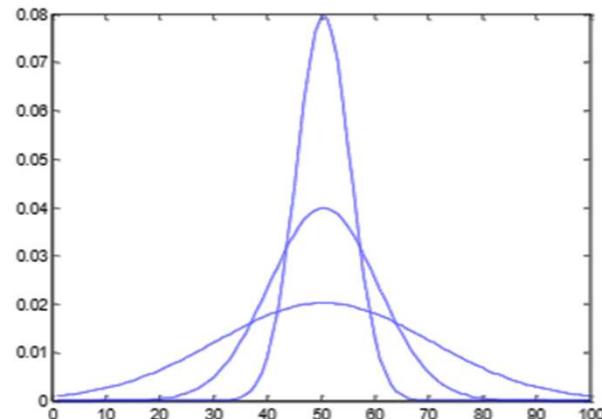


Image courtesy: William Hoff

Scale space

- The space of images created by applying a series of operators of different scales
- Different size Gaussians:

$$G_\sigma(x, y) = e^{-\frac{x^2+y^2}{2\sigma^2}}$$



Lets explore some advanced edge detection techniques

Advanced edge detectors (Marr Hildreth)

- Marr and Hildreth proposed the filter $\nabla^2 G$

- ∇^2 is the Laplacian operator
 - G is 2D gaussian

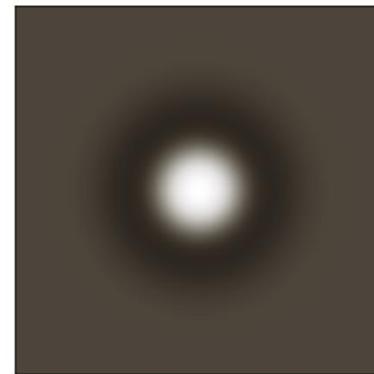
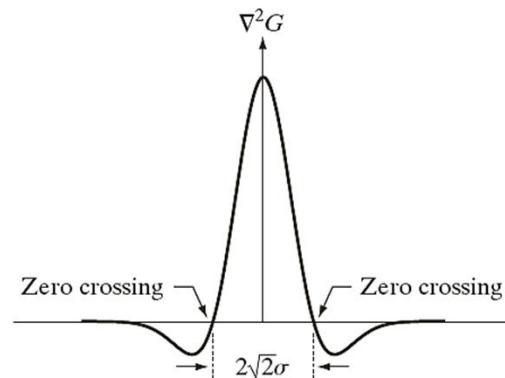
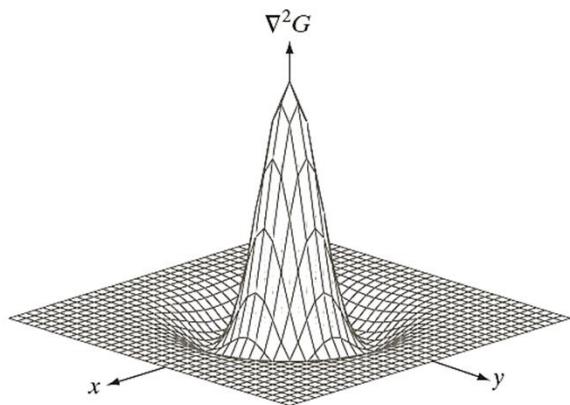
$$G(x, y) = e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$\begin{aligned}\nabla^2 G(x, y) &= \frac{\partial^2 G(x, y)}{\partial x^2} + \frac{\partial^2 G(x, y)}{\partial y^2} = \frac{\partial}{\partial x} \left[\frac{-x}{\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \right] + \frac{\partial}{\partial y} \left[\frac{-y}{\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \right] \\ &= \left[\frac{x^2}{\sigma^4} - \frac{1}{\sigma^2} \right] e^{-\frac{x^2+y^2}{2\sigma^2}} + \left[\frac{y^2}{\sigma^4} - \frac{1}{\sigma^2} \right] e^{-\frac{x^2+y^2}{2\sigma^2}} \\ &= \left[\frac{x^2+y^2-2\sigma^2}{\sigma^4} \right] e^{-\frac{x^2+y^2}{2\sigma^2}}\end{aligned}$$

Laplacian of Gaussian (LoG)!



Advanced edge detectors (Marr Hildreth)



Mexican hat!

a	b
c	d

FIGURE 10.21
 (a) Three-dimensional plot of the negative of the LoG. (b) Negative of the LoG displayed as an image. (c) Cross section of (a) showing zero crossings. (d) 5×5 mask approximation to the shape in (a). The negative of this mask would be used in practice.

0	0	-1	0	0
0	-1	-2	-1	0
-1	-2	16	-2	-1
0	-1	-2	-1	0
0	0	-1	0	0

$$\nabla^2 G(x, y) = \left[\frac{x^2 + y^2 - 2\sigma^2}{\sigma^4} \right] e^{-\frac{x^2 + y^2}{2\sigma^2}}$$

Advanced edge detectors (Marr Hildreth)

- Marr-Hildreth edge detection

$$g(x, y) = [\nabla^2 G(x, y)] * f(x, y)$$

$$g(x, y) = \nabla^2[G(x, y) * f(x, y)]$$

1. Filter the image with a Gaussian low pass filter
2. Compute the Laplacian of the resulting image
3. Find zero crossings

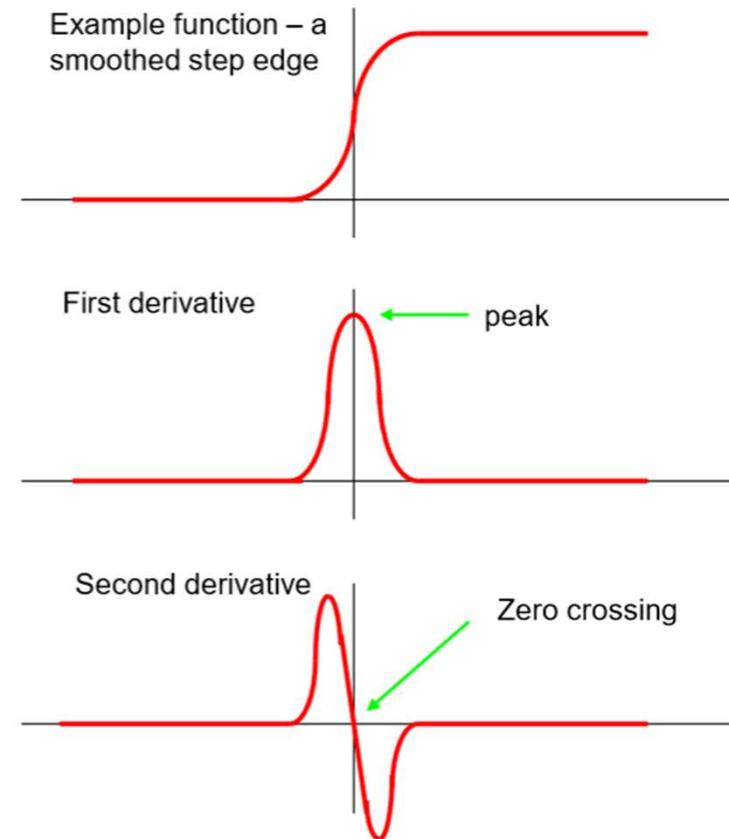
- Advantages:

- Intrinsically apply Gaussian smoothing and avoid ringing noise.
- Gaussian is isotropic (invariant to rotation) so edges of any orientation can be detected.
- Yields better localization and thin edges.



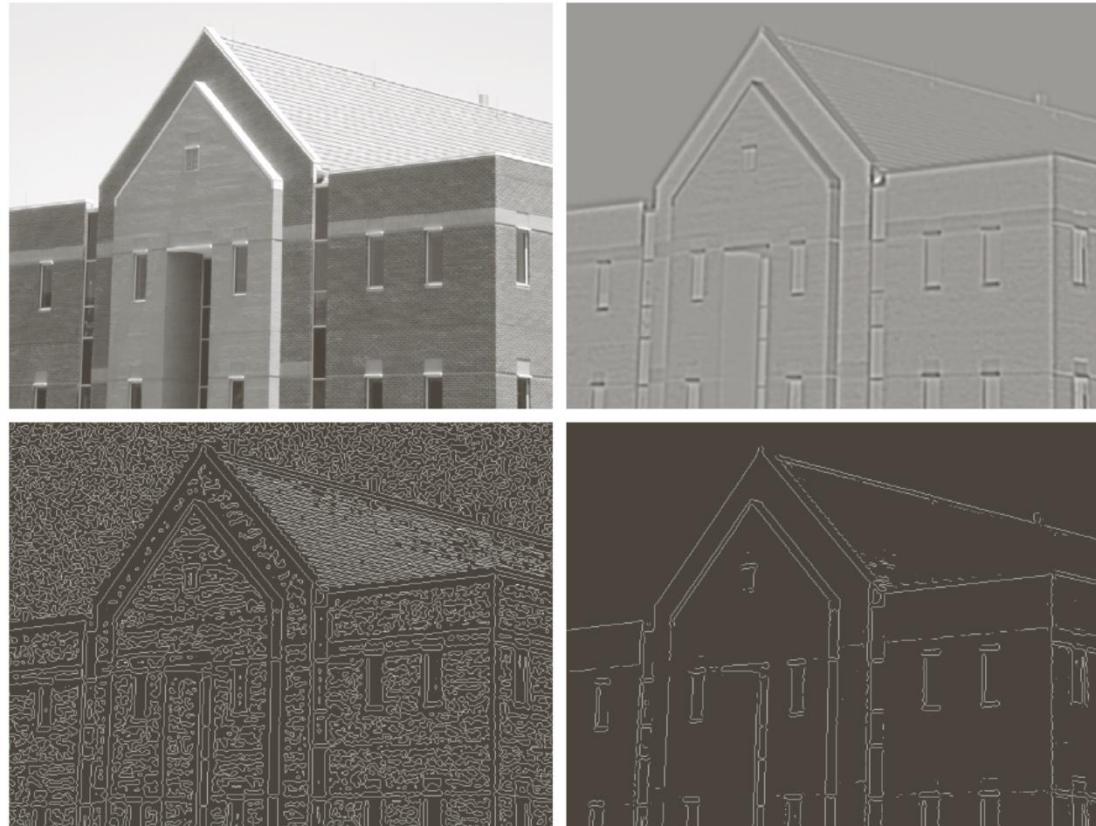
Zero crossings

- Zero crossing indicates an edge
- Peak in first derivative → zeros in second derivative
- How to detect a zero crossing?



Source: William Hoff

Zero crossings



Spaghetti Effect

a b
c d

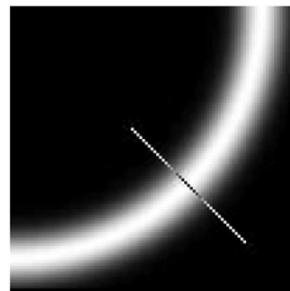
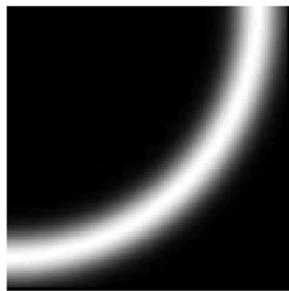
FIGURE 10.22
(a) Original image of size 834×1114 pixels, with intensity values scaled to the range $[0, 1]$. (b) Results of Steps 1 and 2 of the Marr-Hildreth algorithm using $\sigma = 4$ and $n = 25$. (c) Zero crossings of (b) using a threshold of 0 (note the closed-loop edges). (d) Zero crossings found using a threshold equal to 4% of the maximum value of the image in (b). Note the thin edges.

Advanced edge detectors (Canny)

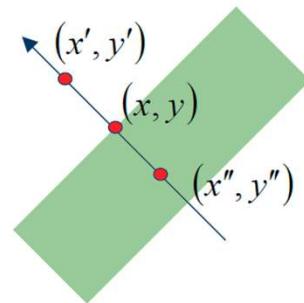
- Most common and widely used
 - Good detection (minimize false edges and missing real edges)
 - Good localization (detected edges close to true edges)
 - Single Response (return only one point for each true edge point)
 - Essence of Canny's work: optimal solution to these three formulation by expressing the three criteria mathematically
 - He found that a good approximation to an optimal operator is the first derivative of the Gaussian in the direction of the gradient
 - Then do non-maxima suppression
 - Then do hysteresis Thresholding
-

Advanced edge detectors (Canny)

- Non-maxima Suppression



Mark points along the curve where the **magnitude** is largest. We can do this by looking for a maximum along a slice normal to the curve (non-maximum suppression).



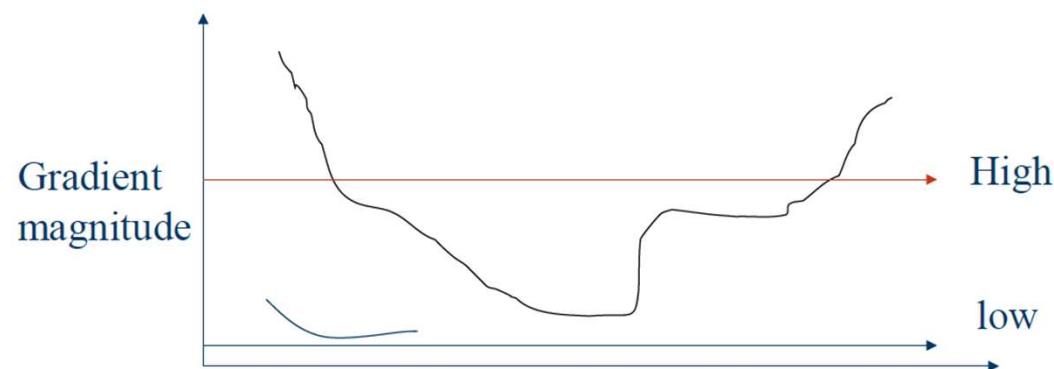
$$M(x, y) = \begin{cases} |\nabla S|(x, y) & \text{if } |\nabla S|(x, y) > |\Delta S|(x', y') \\ & \& |\Delta S|(x, y) > |\Delta S|(x'', y'') \\ 0 & \text{otherwise} \end{cases}$$

x' and x'' are the neighbors of x along normal direction to an edge

Source: Mubarak Shah

Advanced edge detectors (Canny)

- Hysteresis Thresholding
 - If the gradient at a pixel is
 - *above “High”*, declare it as an ‘edge pixel’
 - *below “Low”*, declare it as a “non-edge-pixel”
 - *between “low” and “high”*
 - Consider its neighbors iteratively then declare it an “edge pixel” if it is connected to an ‘edge pixel’ directly or via pixels between “low” and “high”.



Advanced edge detectors (Canny)

1. Smooth the input image with a Gaussian filter
2. Compute the gradient magnitude and angle images
3. Apply non maxima suppression to the gradient magnitude image
 1. Quantize the angle image into four directions
 2. If magnitude is less than one of the two neighbors along its direction, suppress
4. Use hysteresis thresholding and connectivity analysis to detect and link edges



Original



Gradient Magnitude

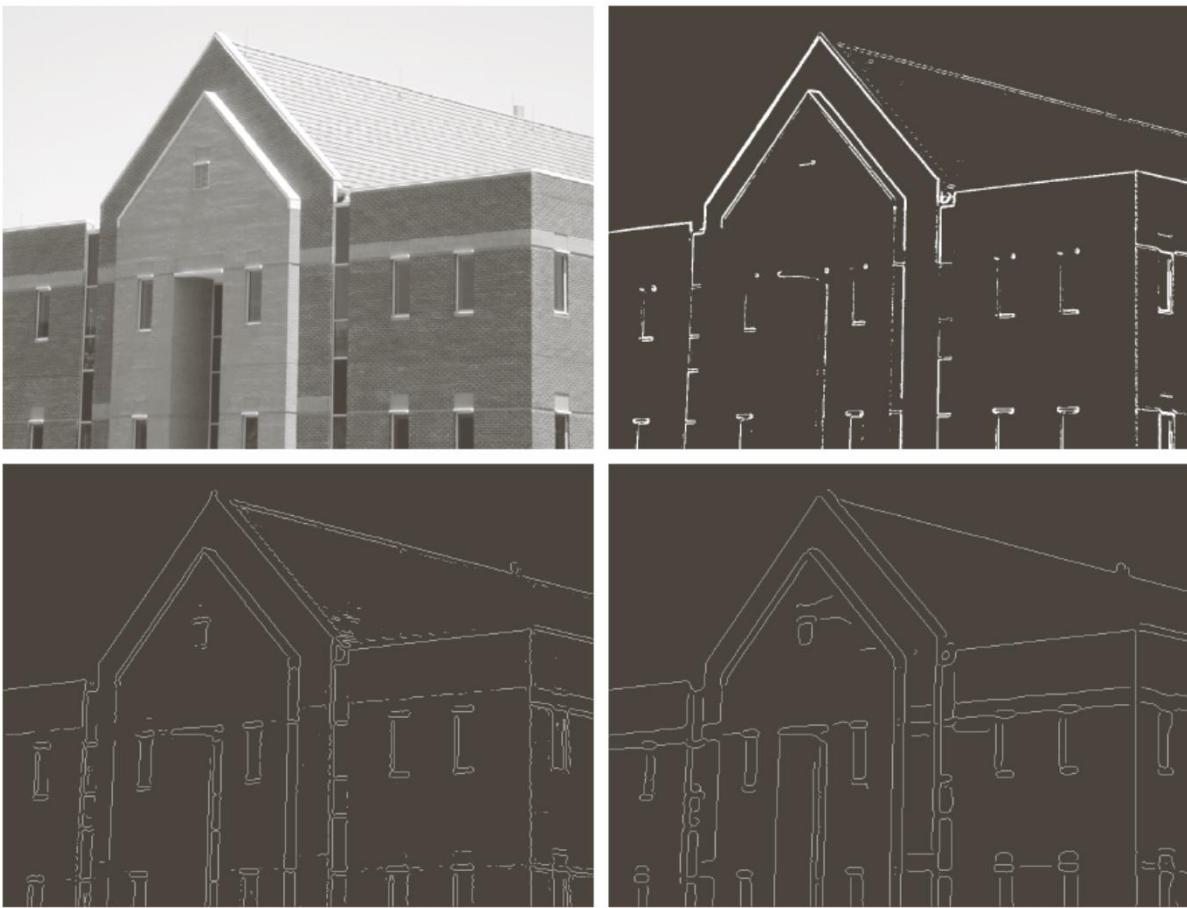


After non-maximal
Suppression



After Hysteresis
Thresholding & Linking

Advanced edge detectors (Canny)



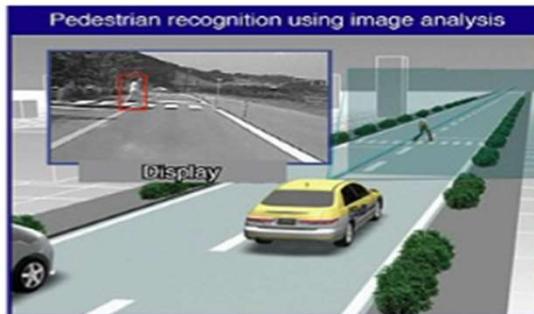
a b
c d

FIGURE 10.25
(a) Original image of size 834×1114 pixels, with intensity values scaled to the range $[0, 1]$.
(b) Thresholded gradient of smoothed image.
(c) Image obtained using the Marr-Hildreth algorithm.
(d) Image obtained using the Canny algorithm. Note the significant improvement of the Canny image compared to the other two.

Lets explore an application of the stuff we learnt!

HOG for person detection

- Applications
 - Multimedia analysis
 - Pedestrian detection for smart cars
 - Visual surveillance, behavior analysis

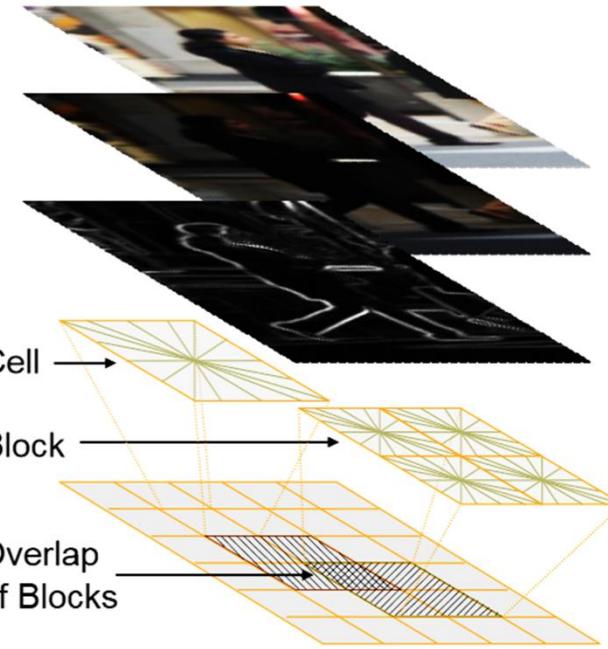
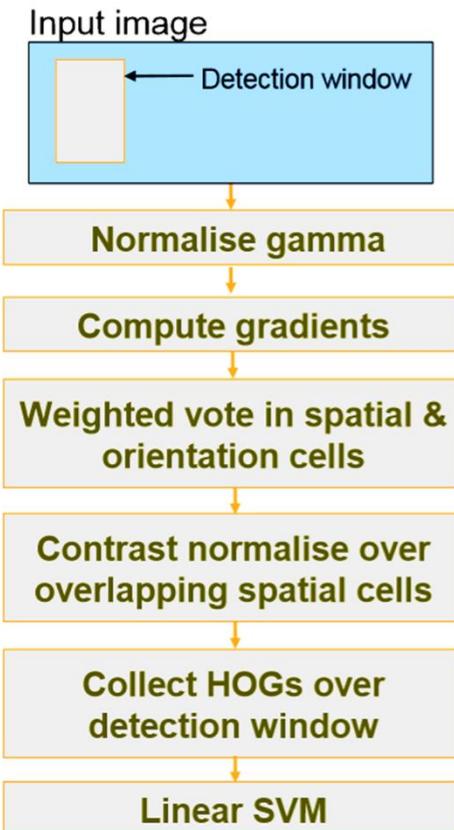


Source: Navneet Dalal

HOG for person detection

- Histogram of oriented gradients
- Proposed by Dalal and Triggs (INRIA Grenoble)
- Extremely popular paper: 20075 citations (todays figure from google scholar)
- Let's understand it in 2 minutes

HOG for person detection

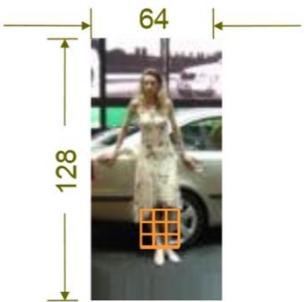
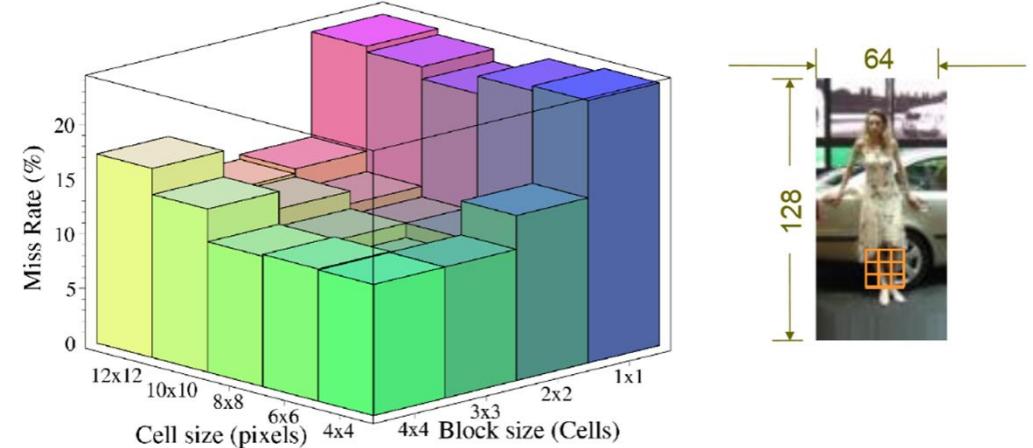
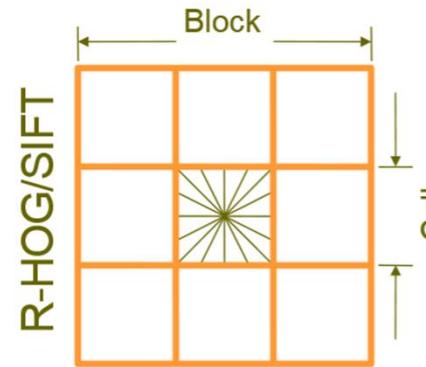


Source: Navneet Dalal

HOG for person detection

Rest is tuning!

- What size of cells? What size of blocks?
- How to normalize?
- How many orientations?
- Which color space?



Source: Navneet Dalal

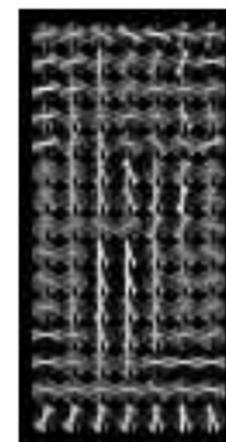
HOG for person detection



(a)



(b)



(c)

HOG detectors cue mainly on silhouette contours (especially the head, shoulders and feet). The most active blocks are centred on the image background just outside the contour. (a) The average gradient image over the training examples (b) A test image. (c) Its computed R-HOG descriptor.

Source: Navneet Dalal

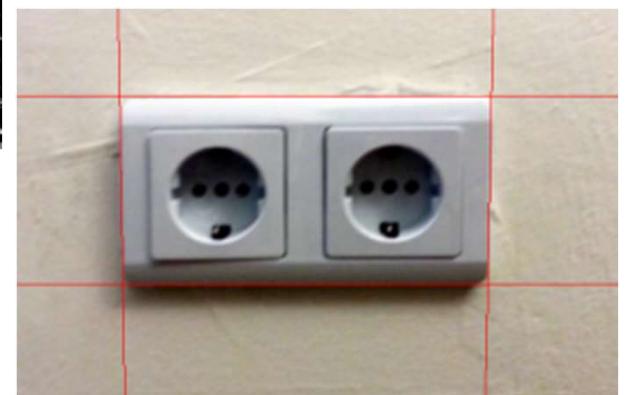
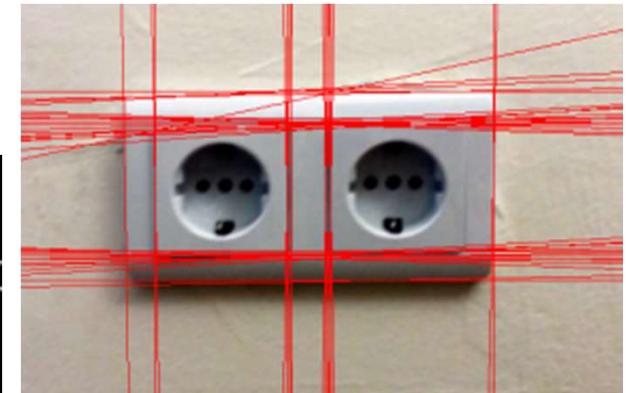
Can we find principle lines or curves from edge images!

Group gradients into Edges or Curves

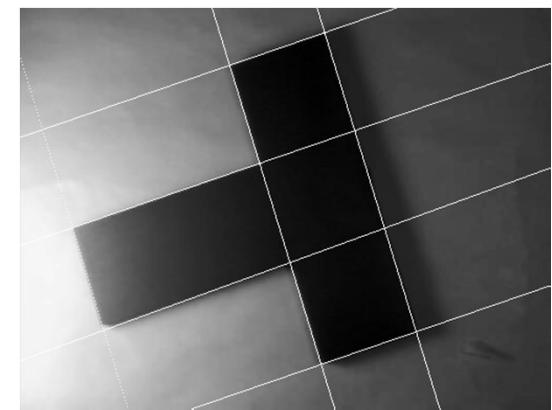
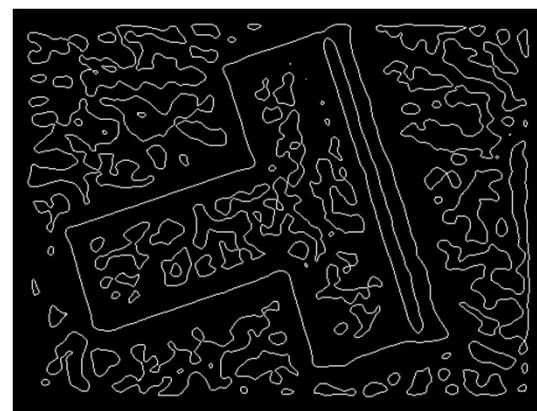
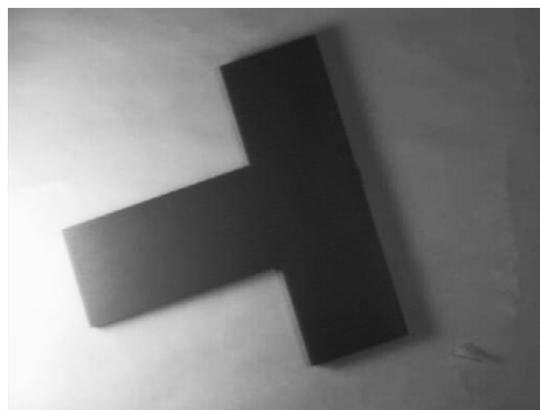
- Local approach
- Global approach



Hough Transform

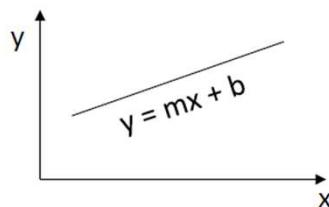


Hough Transform



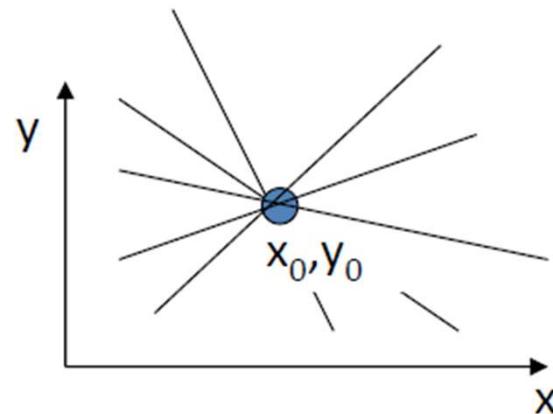
Hough Transform

- Equation of line



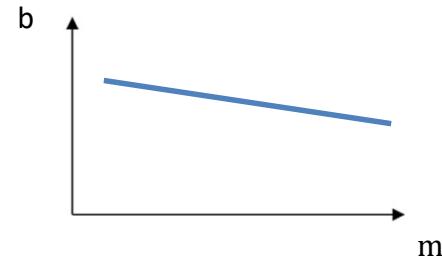
- All lines passing through a point

- In original space



$$b = -x_0 m + y_0$$

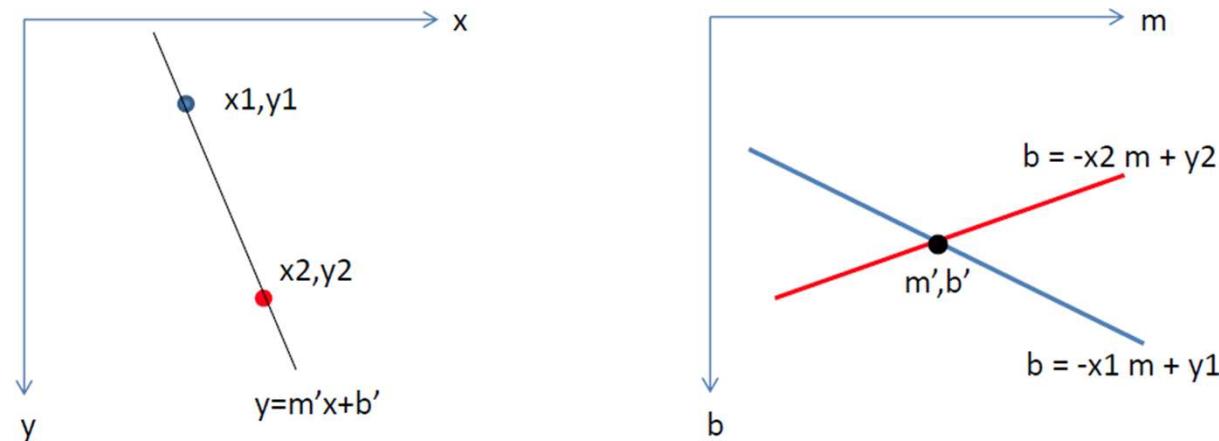
- In parameter space



Source: William Hoff

Hough Transform

- All points on a line in image space, yield lines in parameter space which intersect at a common point
- –This point is the (m, b) of the line in image space

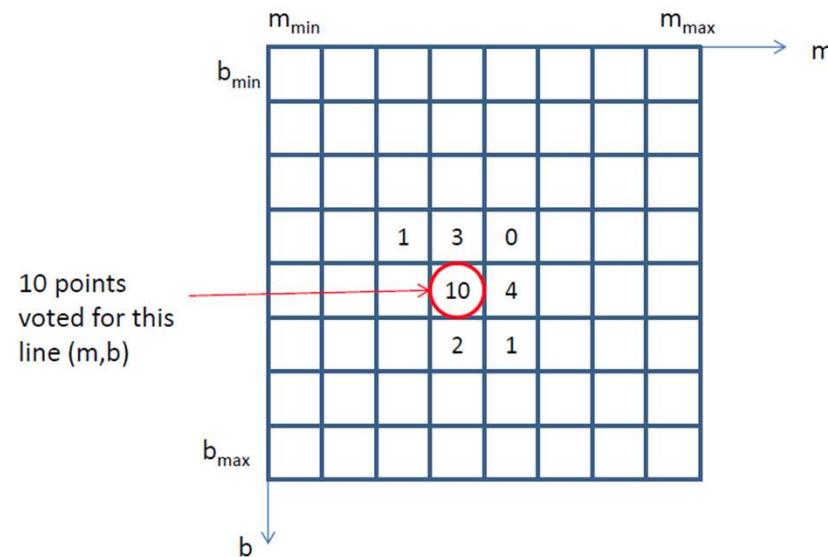


Source: William Hoff

Hough Transform

Algorithm

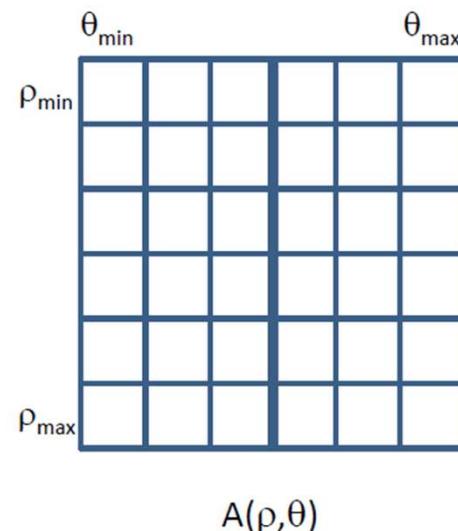
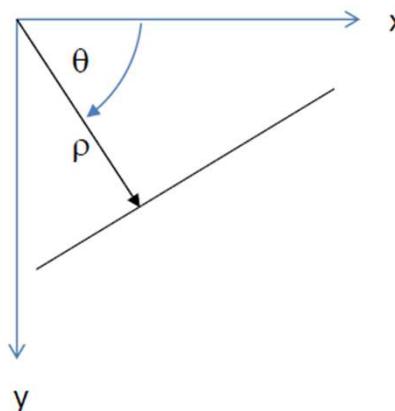
- Initialize an accumulator array $A(m,b)$ to zero
- For each edge element (x,y) , increment all cells that satisfy $b = -x m + y$
- Local maxima in $A(m,b)$ correspond to lines



Source: William Hoff

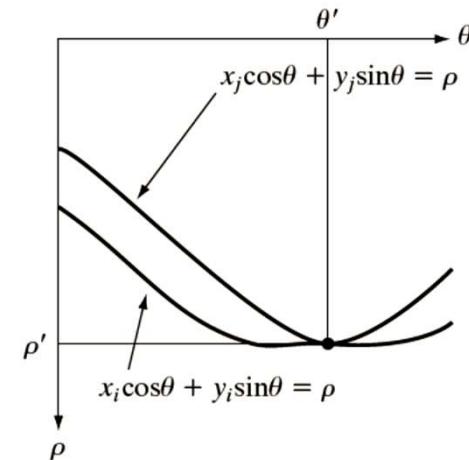
Hough Transform

- $\rho = x \cos \theta + y \sin \theta$
- – Avoids infinite slope
- – Constant resolution



$A(\rho, \theta)$

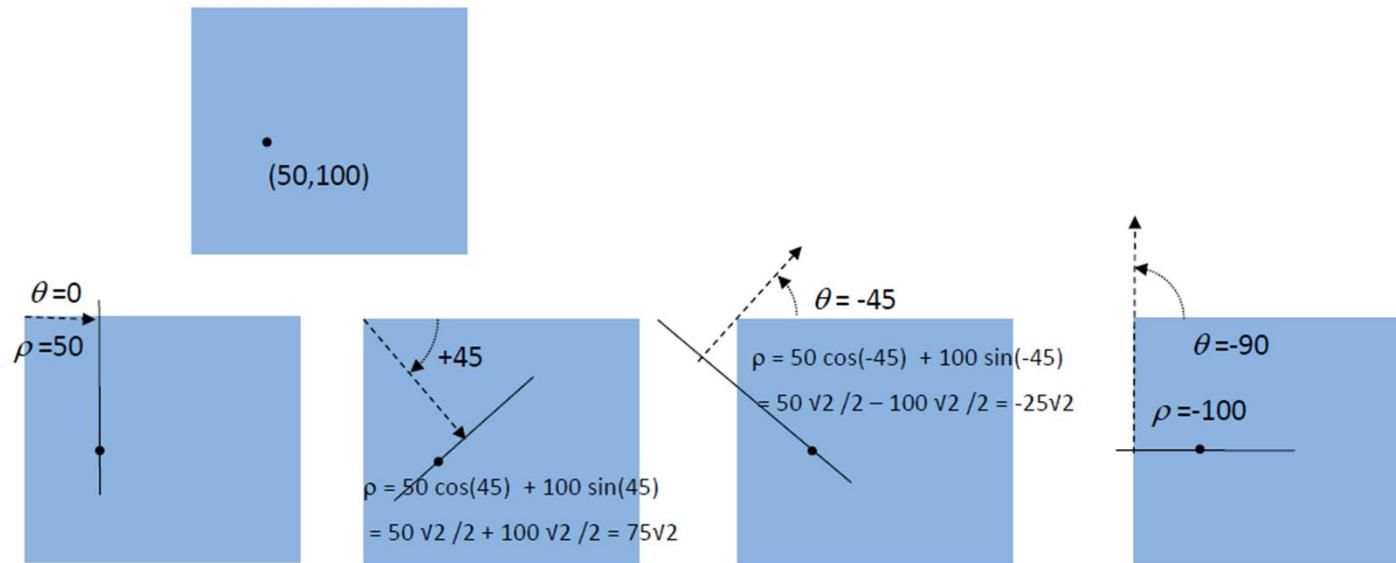
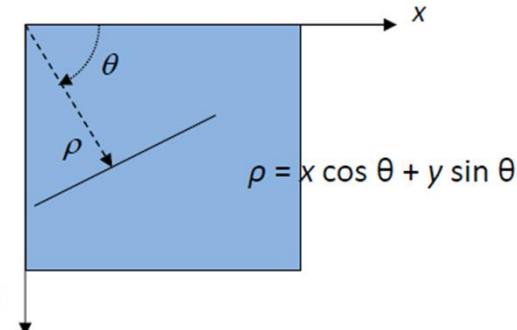
The parameter space transform of a point is a sinusoidal curve



Source: William Hoff

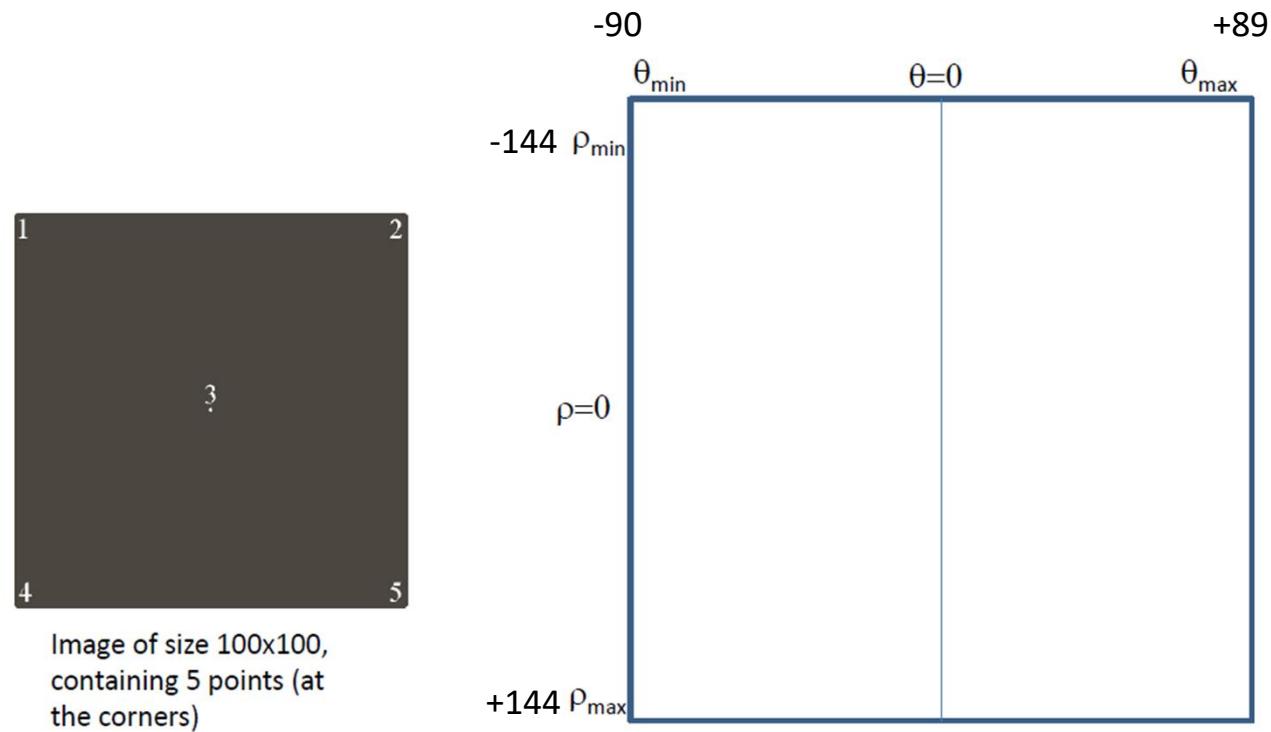
Hough Transform

- Angle, axis conventions
 - angle range is $-90^\circ..+89^\circ$
 - rho range is $-d_{\max}..+d_{\max}$
 - d_{\max} is the largest possible distance
- Example of a point at $(x, y) = (50, 100)$



Source: William Hoff

Hough Transform



Source: William Hoff

Hough Transform

for all x

 for all y

 if edge point at (x,y)

 for all theta

$$\text{rho} = x * \cos(\theta) + y * \sin(\theta)$$

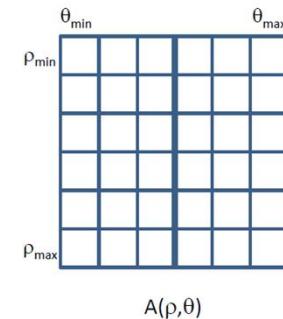
 increment (add 1 to) the cell in A corresponding to (theta,rho)

 end

end

end

end



$A(\rho, \theta)$

Source: William Hoff

Hough Transform

```
clear all
close all
I = imread('gantrycrane.png');

G = rgb2gray(I);
E = edge(G, 'canny');
imshow(E)
[H,theta,rho] = hough(E);
figure, imshow(H, []);

peaks = houghpeaks(H,50,'Threshold',30);

figure, imshow(G, []), hold on

lines =
houghlines(E,theta,rho,peaks,'FillGap',5,'MinLength',15);

for k = 1:length(lines)
    xy = [lines(k).point1; lines(k).point2];
    plot(xy(:,1),xy(:,2),'LineWidth',1,'Color','r');
end
```



Source: William Hoff

Hough Transform

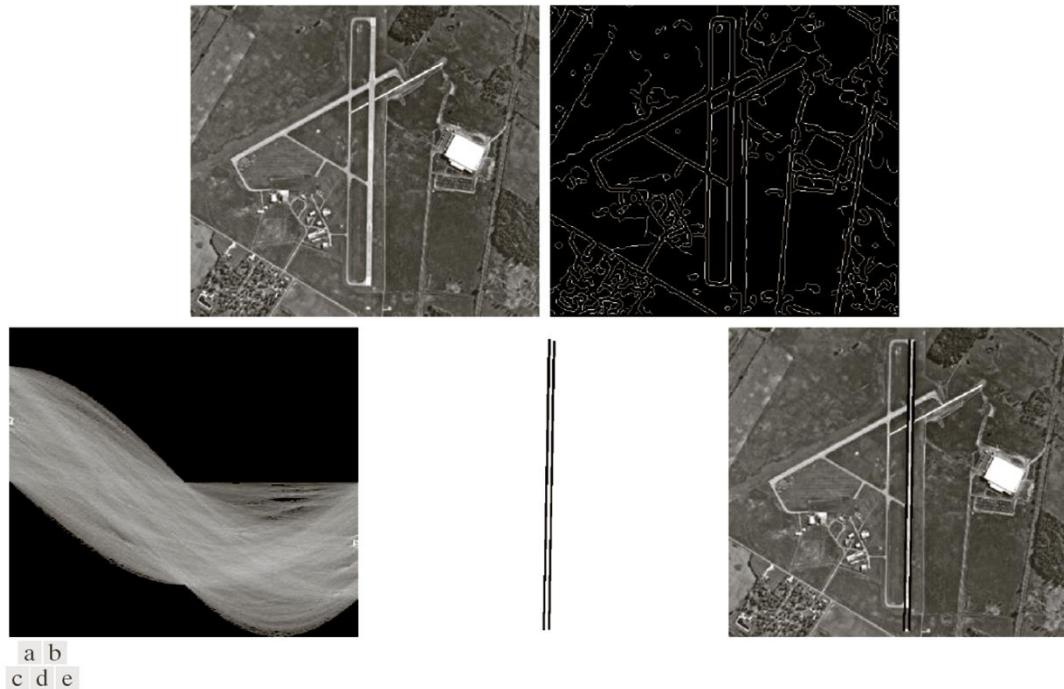


FIGURE 10.34 (a) A 502×564 aerial image of an airport. (b) Edge image obtained using Canny's algorithm. (c) Hough parameter space (the boxes highlight the points associated with long vertical lines). (d) Lines in the image plane corresponding to the points highlighted by the boxes. (e) Lines superimposed on the original image.