

Digital Image Processing (CSE/ECE 478)

Lecture # 21: Image Compression II

Avinash Sharma

Center for Visual Information Technology (CVIT),
IIIT Hyderabad

Lossless compression: Dictionary coding

- Important in presence of recurring patterns
- Static and adaptive
- Static example: a b r a c a d a b r a
- Dynamic dictionary
 - Build during compression after observing the data
 - Rebuild at the decompression step
 - LZW (*Lempel-Ziv-Welch*) is the commonly used algorithm

$A = \{a, b, c, d, r\}$

Code	Entry
000	a
001	b
010	c
011	d
100	r
101	ab
110	ac
111	ad

LZW coding

Lets try to compress the string 'thisisthe'

ASCII codes:

t(116), h(104), i(105),s(115), e(101)

Current	Next	Output	Add to dictionary



LZW coding

ASCII codes:

Lets try to compress the string 'thisisthe'

t(116), h(104), i(105),s(115), e(101)

Current	Next	Output	Add to dictionary
t(116)	h(104)	t(116)	th(256)
h(104)	i(105)	h(104)	hi(257)
i(105)	s(115)	i(105)	is(258)
s(115)	i(105)	s(115)	si(259)
*is(258)			

LZW coding

ASCII codes:

Lets try to compress the string 'thisisthe'

t(116), h(104), i(105),s(115), e(101)

Current	Next	Output	Add to dictionary
t(116)	h(104)	t(116)	th(256)
h(104)	i(105)	h(104)	hi(257)
i(105)	s(115)	i(105)	is(258)
s(115)	i(105)	s(115)	si(259)
*is(258)	t(116)	is(258)	ist(260)
*th(256)	e(101)	th(256)	the(261)
e (101)	-	e(101)	-

LZW decoding

Decompress 116, 104, 105, 115, 258, 256, 101

ASCII codes:

t(116), h(104), i(105),s(115), e(101)

Current	Next	Output	Add to dictionary
116	104	116	116 104 (256)

LZW decoding

Decompress 116, 104, 105, 115, 258, 256, 101

ASCII codes:

t(116), h(104), i(105),s(115), e(101)

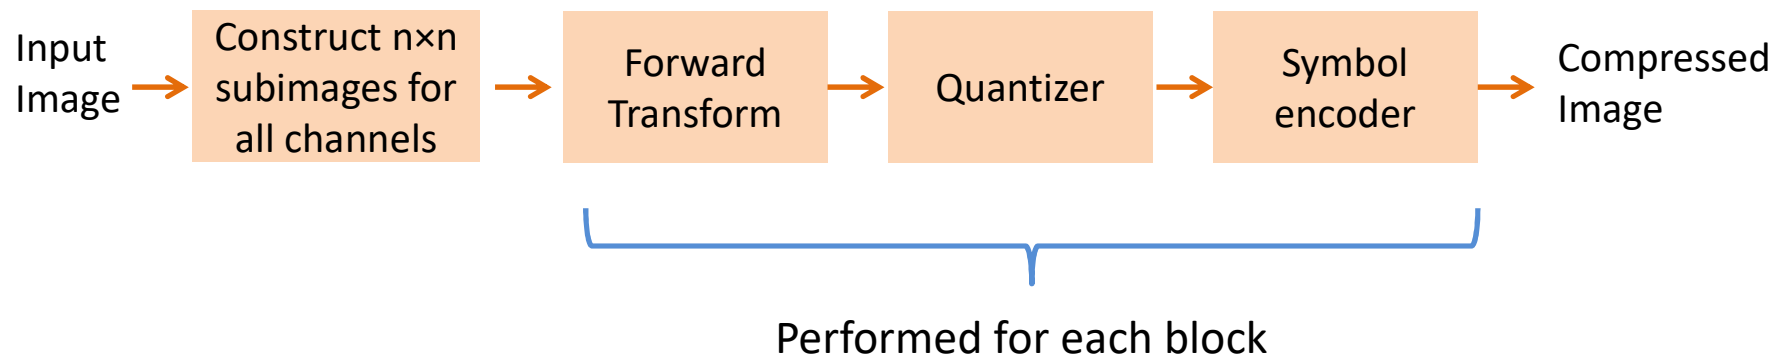
Current	Next	Output	Add to dictionary
116	104	116	116 104 (256)
104	105	104	104 105 (257)
105	115	105	105 115 (258)
115	258	115	115 105 115 (259)
258	256	105 115	105 115 116 104 (260)
256	101	116 104	116 104 101 (261)
101	-	101	-



Lossy compression (with a case study of JPEG)

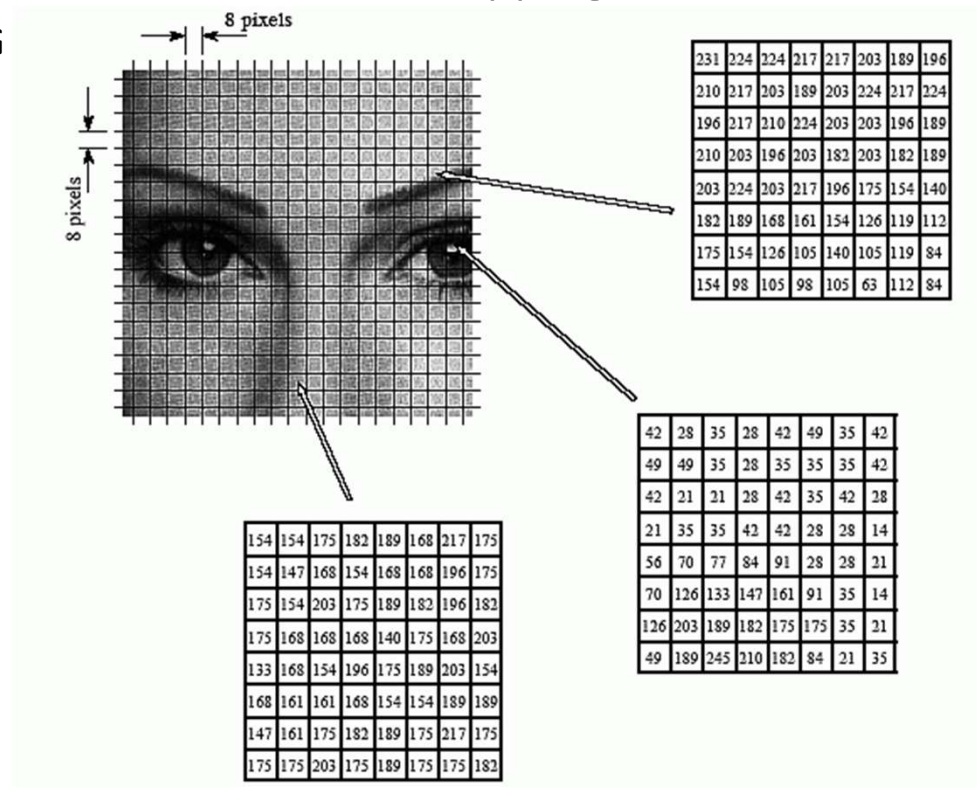


Lossy compression: JPEG



Block transform coding

- Partition the image into small non overlapping $n \times n$ blocks
 - 8×8 blocks in JPEG



Block Transform coding

- Process blocks using 2D transforms
- General forward transform of image g , size $n \times n$:

$$T(u, v) = \sum_{x=0}^{n-1} \sum_{y=0}^{n-1} g(x, y) r(x, y, u, v)$$

- Inverse transform

$$g(x, y) = \sum_{u=0}^{n-1} \sum_{v=0}^{n-1} T(u, v) s(x, y, u, v)$$

- $r(x, y, u, v)$ and $s(x, y, u, v)$ are basis functions or transformation kernels
-


Block Transform coding

$$T(u, v) = \sum_{x=0}^{n-1} \sum_{y=0}^{n-1} g(x, y) r(x, y, u, v)$$


$$g(x, y) = \sum_{u=0}^{n-1} \sum_{v=0}^{n-1} T(u, v) s(x, y, u, v)$$

- $r(x, y, u, v) = e^{-j2\pi(ux+vy)/n}$ and $s(x, y, u, v) = \frac{1}{n^2} e^{j2\pi(ux+vy)/n}$
 Discrete Fourier Transform (DFT)
- $r(x, y, u, v) = s(x, y, u, v) = \frac{1}{n} (-1)^{\sum_{i=0}^{m-1} [b_i(x)p_i(u) + b_i(y)p_i(v)]}$
 Walsh-Hadamard Transform (WHT)
- $r(x, y, u, v) = s(x, y, u, v) = \alpha(u)\alpha(v) \cos\left[\frac{(2x+1)u\pi}{2n}\right] \cos\left[\frac{(2y+1)v\pi}{2n}\right]$
 Discrete Cosine Transform (DCT)

Block Transform coding: which transform to use?



Walsh-Hadamard Transform Basis
for block of size 4×4



Discrete Cosine Transform Basis for
block of size 4×4

Block Transform coding: which transform to use?

Apply transform to each
 8×8 block

Keep highest 50% of the
coefficients in each
block

Reconstruct using the
inverse transform on
each block

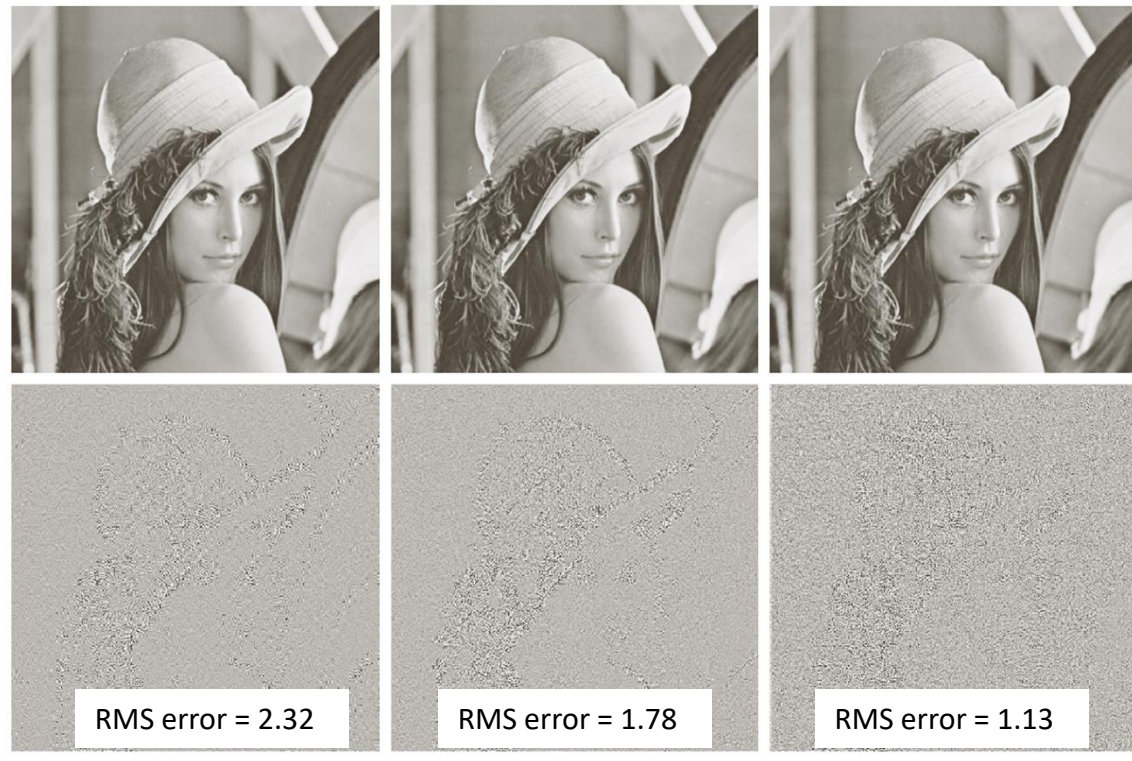


FIGURE 8.24 Approximations of Fig. 8.9(a) using the (a) Fourier, (b) Walsh-Hadamard, and (c) cosine transforms, together with the corresponding scaled error images in (d)–(f).

Block Transform coding: which transform to use?





FIGURE 8.26
Reconstruction
error versus
subimage size.

Block Transform coding: which transform to use?



DFT v/s DST: Issue of *blocking artifact*

Quantization


Need of Quantization

$$\begin{pmatrix} 245 & 239 & 227 & 203 & 174 & 150 & 116 & 92 \\ 229 & 216 & 197 & 172 & 150 & 119 & 85 & 69 \\ 201 & 180 & 164 & 152 & 141 & 102 & 77 & 69 \\ 174 & 153 & 148 & 146 & 140 & 112 & 93 & 91 \\ 161 & 145 & 144 & 146 & 141 & 133 & 120 & 114 \\ 149 & 139 & 143 & 144 & 142 & 139 & 133 & 133 \\ 134 & 128 & 131 & 132 & 134 & 134 & 139 & 137 \\ 119 & 114 & 112 & 111 & 111 & 119 & 131 & 141 \end{pmatrix} \rightarrow \begin{pmatrix} 1 & 7 & 17 & -2 & 1 & 1 & 0 & 0 & 0 \\ 10 & 9 & 16 & -2 & 1 & 0 & 0 & 0 & 0 \\ 4 & 3 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 7 & 5 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 6 & 0 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 6 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \begin{pmatrix} 5.3 & 3.5 \\ 0.6 & 0.3 \\ -1.3 & -2.7 \\ 2.7 & -3.3 \\ 1.1 & -2.5 \\ -1.4 & 2.4 \\ 0.1 & -0.1 \\ -1.8 & 0.2 \end{pmatrix}$$

Three ways to quantize (threshold) a transformed sub-image (block):

- A single global threshold for all blocks
- A different threshold for all blocks (*N-largest coding*, gives constant bit rate)
- Threshold varied as the function of location of each coefficient within the block

Quantization: Threshold Coding



16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	57	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101
72	92	95	98	112	100	103	99

Different coefficients
quantized with different
step-size

$$kc - \frac{c}{2} \leq T(u, v) \leq kc + \frac{c}{2}$$


Finally encode the quantized output!

Quantization (example)

-415	-29	-62	25	55	-20	-1	3		-26	-3	-6	2	2	0	0	0
7	-21	-62	9	11	-7	-6	6		1	-2	-4	0	0	0	0	0
-46	8	77	-25	-30	10	7	-5		-3	1	5	-1	-1	0	0	0
-50	13	35	-15	-9	6	0	3		-4	1	2	-1	0	0	0	0
11	-8	-13	-2	-1	1	-4	1	Q	1	0	0	0	0	0	0	0
-10	1	3	-3	-1	0	2	-1	→	0	0	0	0	0	0	0	0
-4	-1	2	-1	2	-3	1	-2		0	0	0	0	0	0	0	0
-1	-1	-1	-2	-1	-1	0	-1		0	0	0	0	0	0	0	0

16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	57	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101
72	92	95	98	112	100	103	99

Symbol encoding (Zigzag ordering)




0	1	5	6	14	15	27	28
2	4	7	13	16	26	29	42
3	8	12	17	25	30	41	43
9	11	18	24	31	40	44	53
10	19	23	32	39	45	52	54
20	22	33	38	46	51	55	60
21	34	37	47	50	56	59	61
35	36	48	49	57	58	62	63

JPEG uses run length encoding!

Symbol coding example

- Zigzag scan (additional example)



Run length coding




Mean of Block: 185

(0,3) (0,1) (1,1) (0,1) (0,1) (0,-1) (1,1)
 (1,1) (0,1) (1,-3) (0,2) (0,-1) (6,1) (0,-1) (0,-1)
 (1,-1) (14,1) (9,-1) (0,-1) EOF




Lossy compression: JPEG



Let's understand the entire procedure with an example

- Consider a single 8×8 pixel block **B**:

245	239	227	203	174	150	116	92
229	216	197	172	150	119	85	69
201	180	164	152	141	102	77	69
174	153	148	146	140	112	93	91
161	145	144	146	141	133	120	114
149	139	143	144	142	139	133	133
134	128	131	132	134	134	139	137
119	114	112	111	111	119	131	141



- Intensity range $\rightarrow [0 \ 255]$
 - Subtract 127 from each entry and compute 2D DCT
-

Forward transform and quantization

- DCT of image block

$$\hat{\mathbf{B}} = \begin{pmatrix} 118.9 & 187.7 & -17.7 & 16.8 & 14.4 & 2.4 & 5.3 & 3.5 \\ 104.1 & 187.1 & -30.8 & 10.0 & -1.0 & -4.7 & 0.6 & 0.3 \\ 46.3 & 10.4 & 9.1 & -9.0 & -15.7 & 0 & -1.3 & -2.7 \\ 76.8 & -12.1 & -10.7 & -0.2 & -10.4 & 4.8 & 2.7 & -3.3 \\ 6.4 & -15.3 & 1.7 & -1.7 & -1.1 & 2.5 & 1.1 & -2.5 \\ 10.6 & -5.6 & -6.5 & -0.6 & 2.6 & 0.9 & -1.4 & 2.4 \\ 0.4 & -2.3 & 1.2 & -1.7 & 2.3 & -0.5 & 0.1 & -0.1 \\ 3.2 & -0.7 & -0.9 & 2.6 & -1.1 & 1.5 & -1.8 & 0.2 \end{pmatrix}$$

- Quantization and rounding

$$q(\hat{\mathbf{B}}) = \begin{pmatrix} 7 & 17 & -2 & 1 & 1 & 0 & 0 & 0 \\ 9 & 16 & -2 & 1 & 0 & 0 & 0 & 0 \\ 3 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 5 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	57	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101
72	92	95	98	112	100	103	99

More than 75% entries are zero (notice their placement)

Encoding

- Zigzag scan

$$q(\hat{\mathbf{B}}) = \begin{pmatrix} 7 & 17 & -2 & 1 & 1 & 0 & 0 & 0 \\ 9 & 16 & -2 & 1 & 0 & 0 & 0 & 0 \\ 3 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 5 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

0	1	5	6	14	15	27	28
2	4	7	13	16	26	29	42
3	8	12	17	25	30	41	43
9	11	18	24	31	40	44	53
10	19	23	32	39	45	52	54
20	22	33	38	46	51	55	60
21	34	37	47	50	56	59	61
35	36	48	49	57	58	62	63

[7 17 9 3 16 -2 1 -2 1 5 0 -1 1 1 1 0 0 0 0 -1 EOB]

Lets try to reconstruct

Reconstruction: Decoding + Dequantization

[7 17 9 3 16 -2 1 -2 1 5 0 -1 1 1 1 0 0 0 0 -1 EOB]



$$\begin{pmatrix} 7 & 17 & -2 & 1 & 1 & 0 & 0 & 0 \\ 9 & 16 & -2 & 1 & 0 & 0 & 0 & 0 \\ 3 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 5 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	57	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101
72	92	95	98	112	100	103	99


- Dequantization

$$\tilde{\mathbf{B}} = \begin{pmatrix} 112 & 187 & -20 & 16 & 24 & 0 & 0 & 0 \\ 108 & 192 & -28 & 19 & 0 & 0 & 0 & 0 \\ 42 & 13 & 16 & 0 & 0 & 0 & 0 & 0 \\ 70 & -17 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & -22 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$




Decoding

- Compute IDCT and add 127

$$\begin{pmatrix} 254 & 233 & 211 & 197 & 175 & 142 & 110 & 93 \\ 236 & 216 & 194 & 179 & 156 & 120 & 87 & 70 \\ 201 & 184 & 169 & 158 & 138 & 105 & 76 & 62 \\ 166 & 155 & 149 & 147 & 137 & 114 & 94 & 86 \\ 151 & 144 & 143 & 147 & 144 & 130 & 120 & 119 \\ 147 & 140 & 140 & 145 & 143 & 135 & 132 & 137 \\ 133 & 127 & 126 & 130 & 130 & 126 & 131 & 140 \\ 116 & 109 & 109 & 114 & 117 & 118 & 128 & 141 \end{pmatrix}$$


- Compare with original

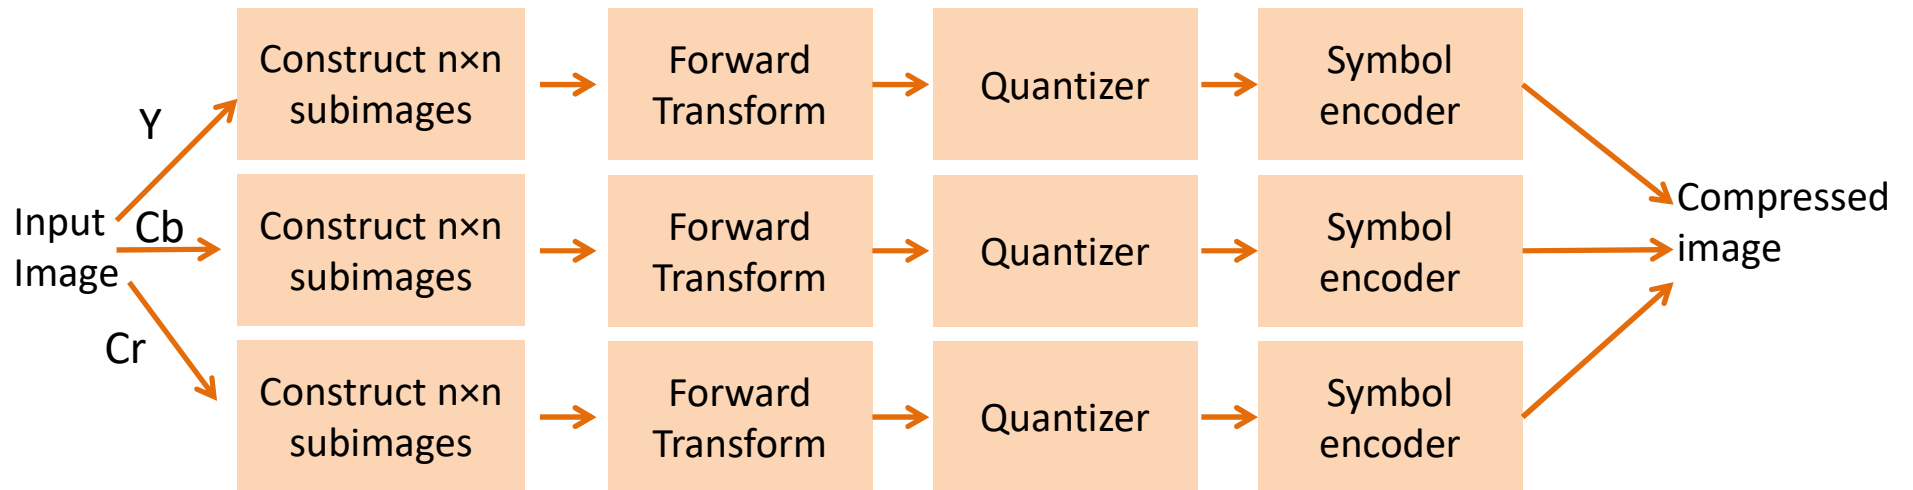
$$\begin{pmatrix} 245 & 239 & 227 & 203 & 174 & 150 & 116 & 92 \\ 229 & 216 & 197 & 172 & 150 & 119 & 85 & 69 \\ 201 & 180 & 164 & 152 & 141 & 102 & 77 & 69 \\ 174 & 153 & 148 & 146 & 140 & 112 & 93 & 91 \\ 161 & 145 & 144 & 146 & 141 & 133 & 120 & 114 \\ 149 & 139 & 143 & 144 & 142 & 139 & 133 & 133 \\ 134 & 128 & 131 & 132 & 134 & 134 & 139 & 137 \\ 119 & 114 & 112 & 111 & 111 & 119 & 131 & 141 \end{pmatrix}$$


Summary JPEG

- Divide into 8×8 subimages
- Compute DCT on each
- Quantize the coefficients
- Order coefficients in zigzag pattern
- Encode 1D sequence using run-length coding and Huffman coding



Color Images



Color Images



Y



Cb



Cr

- Different quantization matrices for chrominance and luminance
- Chroma subsampling (use reduced resolution of chroma channels)



Quantization matrices

16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	57	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101
72	92	95	98	112	100	103	99

Luminance

17	18	24	47	99	99	99	99
18	21	26	66	99	99	99	99
24	26	56	99	99	99	99	99
47	66	99	99	99	99	99	99
99	99	99	99	99	99	99	99
99	99	99	99	99	99	99	99
99	99	99	99	99	99	99	99
99	99	99	99	99	99	99	99

Chrominance

These matrices are scaled for higher compression!